
Evaluating the Role of Great Pre-trained Diffusion Models in Few-shot Phase: Warm-up and Acceleration

Ruofeng Yang^{*1}

Yongcan Li^{*1}

Bo Jiang¹

Cheng Chen²

Shuai Li^{†1,3}

¹Shanghai Jiao Tong University, Shanghai, China

²East China Normal University, Shanghai, China,

³Shanghai Innovation Institute, Shanghai, China,

Abstract

Due to the customized requirements, few-shot diffusion models have attracted much attention. However, only a few works analyze few-shot models, and none involve the fast few-shot optimization process, which is important for quickly responding to users. In this work, we evaluate the role of each operation in the optimization process and prove the convergence guarantee for few-shot diffusion models. A standard operation for the few-shot model is only fine-tuning some key parameters to avoid overfitting the limited target dataset. We first show that this operation is insufficient from empirical and theoretical perspectives. Empirically, we conduct real-world few-shot fine-tuning experiments with underfitting and overfitting bad pre-trained models and show that the results are heavily influenced by these bad models. Theoretically, we also prove that the few-shot phase can not learn the ground-truth parameters and suffers from a small gradient when using a bad pre-trained model. Based on these results, we highlight the importance of a great pre-trained model by showing it can warm up few-shot models and lead to a strongly convex landscape for few-shot diffusion models. As a result, the few-shot model fast converges to the ground-truth parameters. In contrast, we show that with a bad initialization, the pretraining phase requires large optimization steps to converge.

1 INTRODUCTION

Recently, diffusion models, which are trained on large-scale datasets with sufficient training time, have shown impressive performance in different areas such as 2D and 3D gener-

ation [Rombach et al., 2022, Blattmann et al., 2023, Liu et al., 2024]. However, when facing customized requirements, we only have limited data and need to achieve a quick and high-quality response to users. To achieve great performance under such a situation, few-shot diffusion models have received attention [Ruiz et al., 2023, Xiang et al., 2023, Kumari et al., 2023, Moon et al., 2022, Liu et al., 2023]. Few-shot diffusion models only use a limited target dataset (5 – 10 images) and a few optimization steps (fewer than $1k$ steps) to fine-tune a pre-trained model (such as Stable Diffusion (SD) XL, which requires $500k$ optimization steps) and generate samples with the target feature.


Though few-shot diffusion models achieve great performance in applications, only a few works aim to explain the success of few-shot diffusion models [Yang et al., 2024, Chua et al., 2021, Cheng et al., 2025]. Furthermore, they focus on the estimation error and explain why a limited target dataset is enough for few-shot models. However, the fast optimization process of few-shot diffusion models is also important and necessary for quick response to users, and the theoretical guarantee for it is lacking. Hence, the following natural question remains open:

Why do few-shot diffusion models only require a few optimization steps to achieve great performance?

In this work, for the first time, we study the optimization process of few-shot diffusion models, highlight the role of great pre-trained models, and answer the above question by showing that the few-shot phase fast converges to the ground-truth parameters under a suitable condition. Before providing the convergence guarantee, we first analyze the necessary conditions for a great few-shot diffusion model. A standard operation for few-shot diffusion models is to freeze most parameters and only fine-tune some key parameters [Liu et al., 2023, Xiang et al., 2023]. However, we show that this operation is not enough. Empirically, we conduct real-world experiments and show that with bad pre-trained models, the few-shot phase can not generate high-quality images, where overfitting bad pre-trained models suffer from

^{*} These authors contributed equally.

[†] Correspondence to: Shuai Li <shuaili8@sjtu.edu.cn>



Fine-tuning Results based on Great Pre-trained Models (SD3 Medium)



Fine-tuning Results based on *Overfitting* Bad Pre-trained Models (SD3 Medium with 1k overfitting steps)

A *cat* on top of a wooden floor

A *cat* in a chef outfit

A *cat* with a city in the background

A *cat* wearing a yellow shirt

A *cat* in a police outfit

Figure 1: Few-shot fine-tuning results based on great and bad *overfitting* pre-trained Models. The overfitting bad pre-trained model is obtained by training SD3 Medium with 5 *dog* image for 1k steps, which suffers from the memory phenomenon of the overfitting feature in the few-shot phase.

the memory phenomenon (Figure 1) and underfitting bad pre-trained models have a fine-tuning loss gap (Figure 2). Theoretically, we prove that if the pre-trained model is bad, few-shot diffusion models can not learn the ground-truth parameters. Furthermore, the gradient of few-shot diffusion models becomes small when the point is still far away from the minimizer. In other words, few-shot diffusion models require large optimization steps to converge. As a byproduct of the gradient analysis, we also show that the pretraining phase with a bad initialization suffers from a small gradient, which slows down the optimization process.

The above results cannot explain why few-shot diffusion models can achieve great performance with only a few optimization steps. Based on the analysis of bad pre-trained models, we show the importance of great pre-trained models. An intuition is that great pre-trained models provide a warm-up for the few-shot phase and simplify the landscape. Inspired by this intuition, we prove that the few-shot phase with a great pre-trained diffusion model converges to the ground-truth parameters using the gradient descent algorithm with a convergence guarantee.

Combined with the analysis for the pretraining phase, these results explain why few-shot models can use much smaller optimization steps to achieve great performance:

- By providing experiments and counter-examples, we prove the importance of great pre-trained models. Otherwise, the few-shot phase can not learn the ground-truth parameters and suffer slow optimization process.
- We show that with a great pre-trained model, the land-

scape of few-shot diffusion models becomes strongly convex. As a result, few-shot models quickly converge to the ground-truth parameters, and we prove the convergence guarantee for this optimization process.

- As a byproduct of gradient analysis, we prove that the pretraining phase with a bad initialization suffers a small gradient and requires large optimization steps.

Based on our theoretical guarantee and observation, we also provide some practical operation in the real-world few-shot fine-tuning of diffusion models:

• Semantic Alignment Trumps Scale

Latent mismatch in a base model introduces an irreducible error that hinders fine-tuning. When a base model lacks prior knowledge of the target concept, linear adaptations like LoRA cannot salvage it.

• Task-Aware Initialization

Estimate the difficulty of the task. Complex concepts (e.g., faces) significantly increase the cross-term error, demanding strictly higher-quality pre-trained weights to guarantee convergence.

• Trust Fast Adaptation

Great pretraining guarantees local strong convexity and fine-tuning should succeed in very few steps. Prolonged training cannot rescue a bad initialization and early stopping is advised.

Key Takeaway: Choosing a great initialization model is more important than choosing a fine-tuning technique.

2 RELATED WORK

Optimization Guarantee for Diffusion Models. Due to non-convexity, current works either focus on some specific data distribution or use the kernel method to simplify the analysis. For the specific data distributions, a series of works focus on designing algorithms to learn Gaussian Mixture Models [Bruno et al., 2023, Cui and Zdeborová, 2023, Shah et al., 2023, Chen et al., 2024] based on the score matching technique. Han et al. [2024a] focus on a data distribution consisting of two fixed orthogonal vectors. For the general data, Li et al. [2023], Han et al. [2024b], and Bonnaire et al. [2025] simplify the problem to a convex optimization by using a wide 2-layer NN and kernel methods. Then, they use the gradient descent (flow) method to obtain a convergence guarantee. Recently, Zhang et al. [2025b] analyzed the representation learning with a 2-layer ReLU DAE. We note that the above works focus on the pretraining phase. On the contrary, we focus on the few-shot phase, discuss the necessary conditions, and provide the first convergence guarantee for few-shot diffusion models.

Guarantee for Few-shot Diffusion Models. Existing few-shot works primarily bound estimation error by leveraging shared source-target structures [Yang et al., 2024, Cheng et al., 2025]. Only Yang et al. [2024] study the optimization process and provide a closed-form minimizer for the linear subspace distribution with a Gaussian latent. However, the real-world distribution is always multi-modal, and diffusion models usually use optimization algorithms instead of obtaining the closed-form minimizer. In this work, we bridge the gap between few-shot theory and application.

3 PRELIMINARIES

We first introduce the basic knowledge and notation of diffusion models. Let $q_0 \in \mathbb{R}^D$ be the data distribution. The variance preserving (VP) forward process is defined by:

$$dx_t = -x_t dt + \sqrt{2} dB_t, x_0 \sim q_0 \in \mathbb{R}^D,$$

where $\{B_t\}_{t \in [0, T]}$ is a D -dimensional Brownian motion. Let q_t be the density function of x_t and $\{y_t\}_{t \in [0, T]} = \{x_{T-t}\}_{t \in [0, T]}$. To generate samples, diffusion models reverse the forward process and run the reverse process:

$$dy_t = [y_t + 2\nabla \log q_{T-t}(y_t)] dt + \sqrt{2} dB_t.$$

The reverse process requires the score function $\nabla \log q_t(\cdot)$, which contains the data information and can not be exactly calculated. A conceptual way to approximate $\nabla \log q_t(\cdot)$ is to minimize the following score matching (SM) objective function [Song et al., 2020, Karras et al., 2022]:

$$\min_{s \in \text{NN}} \mathcal{L}_{\text{SM}} = \int_{\delta}^T \mathbb{E}_{x_t \sim q_t} \|\nabla \log q_t(x_t) - s(x_t, t)\|_2^2 dt, \quad (1)$$

where NN is a given function class and $\delta > 0$ is the early stopping parameter to avoid a blow-up score. However, $\mathcal{L}_{\text{SM}}$ can not be directly calculated since $\nabla \log q_t(\cdot)$ is unknown for general data. To avoid this problem, Vincent [2011] propose the denoising score matching (DSM) loss based on the conditional score function $\nabla \log q_t(x_t|x_0)$ with an analytical form:

$$\begin{aligned} & \min_{s \in \text{NN}} \mathcal{L}_{\text{DSM}} \\ &= \int_{\delta}^T \mathbb{E}_{x_0} \left[\mathbb{E}_{x_t|x_0} \|\nabla \log q_t(x_t|x_0) - s(x_t, t)\|_2^2 \right] dt, \end{aligned}$$

which is equivalent to $\mathcal{L}_{\text{SM}}$ up to a constant independent of the optimized parameters. Once a forward process is chosen, $q_t(x_t|x_0)$ is determined as $q_t(x_t|x_0) = \mathcal{N}(m_t x_0, \sigma_t^2 I_D)$, and $\nabla \log q_t(x_t|x_0)$ has an analytical form $-(x_t - m_t x_0)/\sigma_t^2$, where $m_t = e^{-t}$, $\sigma_t^2 = 1 - m_t^2$ for VP forward process.

3.1 TWO PHASES OF FEW-SHOT MODELS

Few-shot diffusion models typically operate in two phases: pretraining and few-shot adaptation [Kumari et al., 2023, Moon et al., 2022]. In the pretraining phase, a model is fully optimized on a large source dataset. In the few-shot phase, most parameters are frozen while key components are fine-tuned on limited target data. Following Yang et al. [2024], we model this process by assuming the source (q_s) and target (q_{ta}) distributions share a latent space within linear subspaces.

Assumption 3.1. [Shared Latent in Linear Subspace] Source data x_s and target data x_{ta} have form $x_s = A_s z$ and $x_{ta} = A_{ta} z$ where $A_s, A_{ta} \in \mathbb{R}^{D \times d}$ have orthonormal columns and $z \sim q_z \in \mathbb{R}^d$.

The low-dimensional manifold hypothesis is well-supported empirically [Pope et al., 2021] and widely adopted in theoretical analysis [Chen et al., 2023, Guo et al., 2024]. Similarly, the shared latent assumption is standard in both few-shot diffusion [Yang et al., 2024] and general few-shot learning literature [Du et al., 2020]. With the linear subspace assumption, the score function can be decomposed into (1) a latent score $\nabla \log q_t^{\text{LD}}(\cdot)$ and (2) linear encoder and decoder A_s (A_{ta} for target distribution) [Chen et al., 2023]

$$\nabla \log q_t^s(x) = A_s \nabla \log q_t^{\text{LD}}(A_s^\top x) - (I_D - A_s A_s^\top) x / \sigma_t^2,$$

where $q_t^{\text{LD}}(z') = \int q_t(z'|Z) q_z(z) dZ$ and $q_t(\cdot|z) = \mathcal{N}(m_t z, \sigma_t^2 I_d)$. This decomposition means that the optimization process needs to optimize two parts: the linear encoder and decoder A_s and latent score $\nabla \log q_t^{\text{LD}}(\cdot)$ (represented by ground-truth parameters μ^*). With NN parameters V and μ for linear encoder/decoder and latent score,




Figure 2: The *underfitting* experiments on CelebA64 dataset. Based on the underfitting bad pre-trained models, the few-shot phase can not generate clean face images and suffers from the loss gap.

the objective function for the pretraining phase is

$$\min_{s \in \mathcal{S}_{NN}} \mathcal{L}_{\text{DSM}}^{\text{pre}} = \int_{\delta}^T \mathbb{E}_{x_0 \sim q_s} \left[\mathbb{E}_{x_t | x_0} \left\| \nabla \log q_t^s(x_t | x_0) - s(x_t, t) \right\|_2^2 \right] dt,$$

where $\mathcal{S}_{NN}$ is the function class used in the pretraining phase and has the following form

$$\begin{aligned} \mathcal{S}_{NN} = \{ & \mathbf{s}_{V, \mu}(x, t) = V \mathbf{f}_{\mu}(V^{\top} x, t) / \sigma_t^2 - x / \sigma_t^2 : \\ & V \in \mathbb{R}^{D \times d} \text{ with orthonormal columns,} \\ & \mathbf{f}_{\mu} : \mathbb{R}^d \times [\delta, T] \rightarrow \mathbb{R}^d \text{ a network} \}. \end{aligned}$$

With a pre-trained score function, the diffusion model fine-tunes it with a given target dataset in the few-shot phase. Let $(\hat{V}_s, \hat{\mu})$ be the minimizer of the above pretraining objective function. Since the source and target data share a latent distribution, we freeze the approximated latent score function $\mathbf{f}_{\hat{\mu}}$ and only fine-tune the linear encoder and decoder V_{ta} in the fine-tuning phase:

$$\min_{s \in \mathcal{Q}_{NN}(\hat{\mu})} \mathcal{L}_{\text{DSM}}^{\text{few}} = \int_{\delta}^T \mathbb{E}_{x_0 \sim q_{ta}} \left[\mathbb{E}_{x_t | x_0} \left\| \nabla \log q_t^{ta}(x_t | x_0) - s(x_t, t) \right\|_2^2 \right] dt, \quad (2)$$

where $\mathcal{Q}_{NN}(\mu) = \{ \mathbf{s}_{V, \mu}(x, t) = \frac{1}{\sigma_t^2} V \mathbf{f}_{\mu}(V^{\top} x, t) - \frac{1}{\sigma_t^2} x : V \in \mathbb{R}^{D \times d} \text{ with orthonormal columns} \}.$

Remark 3.2. We denote by I_D the D -dimensional identity matrix and $\mathbf{I}$ the matrix with all elements equal to 1. For a vector $x \in \mathbb{R}^D$, we denote by $\|x\|_2$ the Euclidean norm and $x(i)$ the i -th element. For a matrix $A \in \mathbb{R}^{D \times d}$, we denote by $\|A\|_F$ the Frobenius norm and $A(i, j)$ the (i, j) -th element.

4 INFLUENCE OF BAD PRE-TRAINED MODELS IN FEW-SHOT PHASE

As a start, we conduct experiments to show the influence of bad pre-trained models in the few-shot phase. The bad

		Great	Bad1k	Bad4k
SD1.4	Clip-T	0.326	0.324	0.322
	Pickscore	21.70	21.70	21.67
SD3-M	Clip-T	0.324	0.309	0.236
	Pickscore	22.11	20.52	18.52

Table 1: Results for overfitting bad pretraining.

pre-trained models can be roughly divided into overfitting and underfitting, where the former suffers from low diversity and the latter does not learn the basic information.

Bad Overfitting Pretraining: Memory Phenomenon. To simulate the memory phenomenon, we deliberately overfit standard Stable Diffusion (SD) models using a single prompt ("a photo of a dog") and 5 images. This yields two degraded baselines: Bad1k and Bad4k, trained for 1k and 4k steps, respectively. We observe that increased overfitting steps directly correlate with a significant reduction in generation diversity. In the few-shot phase, following Dreambooth [Ruiz et al., 2023], we fine-tune the pre-trained model with Dreambooth training dataset. We evaluate performance using CLIP-T and PickScore calculated on 3k images generated from test prompts. To match Section 3.1, we only fine-tune the first and last 3 blocks of NN (Details in Section B). As shown in Figure 1, if the pre-trained models overfit a dog feature, the few-shot phase suffers from the memory phenomenon and can not generate images with the target feature (cat). Table 1 also shows that as the pretrained models become worse, the metric is worse.

Bad Underfitting Pretraining: Few-shot Loss Gap. Following Yang et al. [2024], we conduct experiments on CelebA64 to show the influence of underfitting pre-trained models. We first train two basic models with different hairstyles. The facial features generated by the basic model trained with 50 epochs (Model-50) are distorted, while Model-200 (with converged loss) can generate clear faces. Hence, we call Model-50 an underfitting model and Model-200 a great pre-trained model.

After fine-tuning on bald targets, Model-50 fails to generate




Figure 3: Low-diversity OF pretraining produces near-duplicates of the training subset. Each panel: top row = generated samples, bottom row = nearest source-training image, $d = \text{LPIPS distance}$. (a) Bald OF200 (200-image source): $d \approx 0.001$. (b) Bald M200 (full 6400, matched budget): $d \approx 0.06 (\approx 100\times \text{larger})$. (c) Eyeglasses OF200_g: same $d \approx 0.001$ signature (cross-concept replication).




Figure 4: Quantitative bald rate across pretraining checkpoints. M-series (blue) shows monotone improvement with pretraining maturity (3.7% $\rightarrow$ 73.6%); OF-series (red) at compute-matched budgets shows -11.6 pp and -8.4 pp deficits relative to M200 / M400.

coherent facial structures (Figure 2 a), manifesting as a persistent loss gap compared to Model-200 (Figure 2 b). This suggests that fundamental concepts missing from pre-training cannot be recovered during few-shot adaptation, leading to suboptimal local minima.

Empirical Evidence with Standard Pretraining. Compared to simulate overfitting starting from good SD3 checkpoint, a more natural way is to take under-trained checkpoints from a standard pretraining trajectory. Hence, we pre-train a series of diffusion models on CelebA64 (6400 pretraining images) and take early-stage checkpoints as the underfitting bad pretraining models. For the overfitting bad models, we pretrain models on a smaller subset of those 6400 images at the same optimizer-step budget.

More specifically, we train diffusion models on 6400 images for 400 epochs and take five checkpoints with difference epochs (where MN denotes the checkpoint after N pre-training epochs) as checkpoints from a single 6400-image trajectory. OF200 and OF500 are checkpoints trained at matched optimizer budgets on 200 and 500 image source subsets. The fine-tuning phase uses 2000 steps on 10 target images with three independent target seeds, identical across all checkpoints.

Metrics.

Bald rate: Fraction of 3000 generated samples scored $P(\text{Bald}) > 0.5$ by a 3-seed CelebA-trained attribute-classifier

ensemble, which is used to measure adaptation success to the target data.

NN-LPIPS / near-copy rate ($\text{NC}@ \tau$): nearest-source-image LPIPS distance for each generated sample, with $\text{NC}@ \tau$ defined as the fraction of samples whose NN-LPIPS falls below a pre-calibrated threshold $\tau = 0.0730$ (set on real bald hold-out faces before any generated sample was inspected), which is used to measure pretrained-stage memorization.

Bad overfitting pretraining suffers from memorization.

Using NN-LPIPS nearest-neighbor checks, the low-diversity pretrained checkpoints show clear pretrained-stage memorization: OF200 has $\text{NC}@ \tau = 94.93\%$ versus 0.63% for matched-budget M200, and as shown in Figure 3, OF200 and OF500 pretrained checkpoints generate near-duplicates of the small training subset.

Results for few-shot phase with bad pretraining.

After identical 10-shot fine-tuning, both undertrained and low-diversity / high-exposure checkpoints achieve lower Bald rates than the better-pretrained M200 / M400 checkpoints. As shown in Figure 4 and Figure 5, the generated images are visibly worse under bad pretraining than under good pretraining.

5 BAD PRETRAINING PREVENTS FEW-SHOT LEARNING

An intuitive idea is that the few-shot phase will quickly converge to the global minimizer since it only optimizes fewer parameters than the pretraining phase. However, we show that if the pre-trained model is bad (the criteria of great pretrained models is provided in Section 6.1), the few-shot phase can not learn A_{ta} , suffers a constant error (few-shot loss gap), and has a small gradient.

Before providing our results, we first make some assumptions on the latent distribution q_z . As discussed in Sec.2, Yang et al. [2024] only adopt a Gaussian latent, which can not reflect the multi-modal property of the real-world dataset. On the contrary, we assume the latent distribution is a mixture of Gaussians [Shah et al., 2023], which is multi-modal and has a nonlinear score function.

Assumption 5.1. The source and target data lie on linear




Figure 5: Visual fine-tuning results across pretraining checkpoints. All five fine-tuned on the same 10 bald target images; only pretraining differs. Each panel: $4 \times 4 = 16$ LPIPS-farthest-first diverse samples. (a)-(c) M-series dose response (bald rate 3.7% $\rightarrow$ 73.6%). (d), (e) OF checkpoints at compute-matched budgets show -11.6 pp and -8.4 pp adaptation deficits relative to M200 / M400.

subspaces $x_s = A_s z$ and $x_{ta} = A_{ta} z$, defined by distinct orthonormal mixing matrices $A_s, A_{ta} \in \mathbb{R}^{D \times d}$. The latent variable $z \in \mathbb{R}^d$ follows a Gaussian mixture distribution with density $q(z) = \sum_{k=1}^K \pi_k \mathcal{N}(z; \mu_k, \Sigma_k)$.

Remark 5.2. While Assumption 5.1 defines a general GMM, we adopt a canonical symmetric 2-mode setting ($D = 2, d = 1$) to explicitly quantify the failure of bad pretraining and, in Section 6, to derive closed-form convergence rates for great pretraining. Crucially, we prove that the geometric guarantees extend to general K -mode GMMs in arbitrary ambient dimensions D (Appendix F), while the generalization to high-dimensional latent spaces ($d > 1$) is further supported by intuition and simulations in Tables 2 and 4.

Let $\mu_t^* = \mu^* \exp(-t)$. After assuming 2-modal Gaussian Mixture latent, the ground truth latent score $\nabla \log q_t^{\text{LD}}(\cdot)$ has a closed form, which leads to the following score in the full space:

$$\nabla \log q_t(x) = \frac{A \tanh(\mu_t^{*\top} A^\top x) \mu_t^* - A A^\top x - (I_D - A A^\top) x / \sigma_t^2}{\sigma_t^2}. \quad (3)$$

Inspired by the above formulation, we use the following network $\mathbf{f}_\mu(z, t)$ to approximate latent score:

$$\mathbf{f}_\mu(z, t) = \sigma_t^2 \tanh(\mu_t^\top z) \mu_t + (1 - \sigma_t^2) z,$$

and $V_{ta} = [v_{ta}, v_{ta}]^\top$. After determining $\mathcal{S}_{\text{NN}}$ and $\mathcal{Q}_{\text{NN}}$, the few-shot diffusion models can first optimize $\mathcal{L}_{\text{DSM}}^{\text{pre}}$ and fine-tune the pre-trained score with $\mathcal{L}_{\text{DSM}}^{\text{few}}$ with the target dataset. We also define the score matching objective function for the few-shot phase, which is used in the analysis:

$$\min_{s \in \mathcal{Q}_{\text{NN}}(\hat{\mu})} \mathcal{L}_{\text{SM}}^{\text{few}} = \int_{\delta}^T \mathbb{E}_{x_t} \left\| \nabla \log q_t^{ta}(x_t) - s(x_t, t) \right\|_2^2 dt,$$

where $x_t \sim q_t^{ta}$. We note that $\mathcal{L}_{\text{SM}}^{\text{few}}$ and $\mathcal{L}_{\text{DSM}}^{\text{few}}$ are equivalent up to a constant independent of all optimized parameters [Vincent, 2011], which indicates the optimization landscape is the same (Fig. 6). Since the score $\nabla \log q_t(\cdot)$ under Theorem 3.1 and 5.1 has an analytical form (Eq. 3), we focus on

the score matching (SM) objective function. We note that when considering the convergence guarantee of the pretraining phase, Li et al. [2023] also adopt the score matching objective $\mathcal{L}_{\text{SM}}^{\text{pre}}$.

Notation. The learnable parameters can be divided into linear encoder/decoder and latent score. For linear encoder/decoder, A_s, V_s , and $\hat{V}_s$ are the ground-truth parameters, learnable NN parameters, and empirical minimizer for the source data, respectively (A_{ta}, V_{ta} and $\hat{V}_{ta}$ for target data). For the shared latent score $\nabla \log q_t^{\text{LD}}$, we use μ^*, μ and $\hat{\mu}$ as the ground-truth, NN and minimizer parameters.

Regarding objectives, we focus on Score Matching ($\mathcal{L}_{\text{SM}}$) over Denoising Score Matching ($\mathcal{L}_{\text{DSM}}$) due to their shared landscape. Superscripts denote phases: pre (pretraining), few (trainable V_{ta}), and few - all (trainable V_{ta}, μ). Our analysis primarily concerns the influence of a pretrained $\hat{\mu}$ under the few setting. Since A is time-independent, we fix $t \in [\delta, T]$ and abbreviate the main objective $\mathcal{L}_{\text{SM},t}^{\text{few}}$ as $\mathcal{L}$. Otherwise, we clearly show the subscript and superscript.

5.1 RESULTS WITH A BAD PRETRAINING

For convenience, we call the pre-trained model great when $\hat{\mu} = \mu^*$. Otherwise, we call the pre-trained model bad *. We note that the underfitting and overfitting pre-trained models satisfy this definition, where the former do not learn basic concepts and the latter overfit to some specific features. Let $\hat{V}_{ta}$ be the solution of $\partial \mathcal{L} / \partial V_{ta} = 0$. In this part, we show that with a bad pre-trained model, $\|\hat{V}_{ta} \hat{V}_{ta}^\top - A_{ta} A_{ta}^\top\|_F$ is not equal to 0, which indicates the few-shot phase can not learn ground-truth subspace parameters and suffers a few-shot loss gap.

Lemma 5.3. Assume Theorem 3.1 and 5.1. If $\hat{\mu} \neq \mu^*$, with $V_{ta} V_{ta}^\top = A_{ta} A_{ta}^\top$, $\partial \mathcal{L} / \partial V_{ta} \neq 0$.

Since the source dataset are limited, $\|\hat{\mu} - \mu^\|$ is smaller than a small constant ϵ_{pre} instead of equal to 0 for a great pre-trained model. In Section 6.1, we discuss the influence of limited data and imperfect learning.




Figure 6: The Landscape for $\mathcal{L}_{SM}$ and $\mathcal{L}_{DSM}$. Since the landscape of $\mathcal{L}^{\text{few}}$ can be viewed as a slice of $\mathcal{L}^{\text{few-all}}$ at $\hat{\mu}$, we present the landscape of $\mathcal{L}^{\text{few-all}}$.

This lemma indicates that $\|\hat{V}_{ta}\hat{V}_{ta}^\top - A_{ta}A_{ta}^\top\|_F \neq 0$ if $\hat{\mu} \neq \mu^*$, which explain the few-shot loss gap in Fig.2. Then, we discuss the influence of $|\hat{\mu} - \mu^*|$ by using a simplified example $\hat{\mu} = 0$ and $\mu^* \neq 0$.

Theorem 5.4. Assume Theorem 3.1 and 5.1 hold. Let μ_1^* and μ_2^* be the two parameters to generate different latent distributions. Given a bad pre-trained model with $\hat{\mu} = 0$, if $|\mu_1^* - \hat{\mu}| > |\mu_2^* - \hat{\mu}|$, then

$$\|\hat{V}_{ta,1}\hat{V}_{ta,1}^\top - A_{ta}A_{ta}^\top\|_F > \|\hat{V}_{ta,2}\hat{V}_{ta,2}^\top - A_{ta}A_{ta}^\top\|_F,$$

where $\hat{V}_{ta,i}$ is the solution corresponds to μ_i^* , $i \in \{1, 2\}$.

This result shows that with a worse pre-trained model, the solution of the few-shot phase becomes worse. Hence, a great pre-trained model is necessary for the few-shot phase. Before providing positive results, we further prove that with a bad pre-trained model, another fully fine-tuning method for few-shot models also has a bad performance and suffers from a small gradient.

Fully Fine-tuning Method and Gradient Analysis. Though many empirical works only fine-tune key parameters, a few work [Ruiz et al., 2023] still optimize all parameters $\min_{s \in S_{NN}} \mathcal{L}_{SM}^{\text{few-all}}$ in the few-shot phase with initialization $(\hat{V}_s, \hat{\mu})$, where $\mathcal{L}_{SM}^{\text{few-all}}$ is the same with Eq. 1. In the following theorem, we show that with a bad pre-trained model, $\partial \mathcal{L}_{SM,t}^{\text{few-all}} / \partial \mu_t$ is small when the point is far away from the global minimizer.

Theorem 5.5. Assume Theorem 3.1 and 5.1 holds. For a fixed t , if $\mu_t \in (-\epsilon, \epsilon)$, we have that

$$\partial \mathcal{L}_{SM,t}^{\text{few-all}} / \partial \mu_t \leq 4\epsilon A_{ta}^\top V_{ta} \sqrt{(1 + \mu_t^{*2}) V_{ta}^\top V_{ta} \sqrt{C_1}} + O(\epsilon^{\frac{3}{2}}),$$

where C_1 is a small constant determined by V_{ta} , A_{ta} and μ^* (Details in Eq. 10).

The above result indicates that if $\hat{\mu}_t \in (-\epsilon, \epsilon)$, the gradient is small. Then, if μ^* is a positive constant larger than ϵ , the few-shot phase requires large optimization steps to get rid of the bad pretraining phase. We use a toy example to show the scale of the gradient, which is much smaller than ϵ .

Example 5.6. Considering $A_s = [0.1, 0.1]^\top$, $A_{ta} = [0.12, 0.12]^\top$, and $\mu^* = 4$. With a fixed $t = 2$ and $V_{ta} = [0.1, 0.1]$ (close to the A_{ta}), $\partial \mathcal{L}_{SM,t}^{\text{few-all}} / \partial \mu_t \leq 1 \times 10^{-5}$ when $\mu_t \in (-0.12, 0.12)$.

Remark 5.7 (Pretraining phase). Since the fully fine-tuning objective function is the same as the pretraining one (only different in the dataset), this result can also explain why the pretraining phase requires large optimization steps. More specifically, since the pretraining phase does not have the prior information of μ^* , it is possible to initialize μ around 0, which leads to a small gradient.

Remark 5.8 (Practical Operation). Sec. 5 supports our first operating principle. Inadequately aligned pre-trained models can introduce irreducible errors during few-shot phase, thus affecting the fine-tuning results.

6 GREAT PRETRAINING: WARM-UP AND ACCELERATION OPTIMIZATION

2-mode GMM analysis. To derive explicit convergence rates and quantify the interplay between pretraining error and data complexity, we first consider a symmetric 2-mode GMM latent distribution with dimension $d = 1$ and symmetric mixing matrices in ambient dimension $D = 2$.

A significant advantage of the few-shot phase is that it can use the information of a well-trained score as the prior (such as latent information μ and data structure), which provides a warm-up for the few-shot phase. Based on this intuition, we show that few-shot models enjoy a simplified landscape and quickly converge to ground-truth parameters with a great pre-trained model.

To achieve this goal, we prove that the landscape of few-shot phase is strongly convex with great pretraining. As a start, we first show $\partial^2 \mathcal{L} / \partial V_{ta}^2 = 2(N + M)$ consists two parts: the first squared term N and the second cross term M (we ignore (x_t, t) and $\mathbb{E}_{x_t \sim q_t^{ta}}$ for clarity):

$$\begin{aligned} \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) &:= N, \\ \left(\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) &:= M. \end{aligned}$$

We know that the squared term N is a semi-positive definite matrix. However, due to the influence of the cross term, we determine a more precise lower bound for each element of N , as shown in the following lemma (In the following two lemmas, we ignore the ta index of v_{ta} and V_{ta}).

Lemma 6.1. [Squared Term] Assume Theorem 3.1 and 5.1 holds and the latent parameter $\hat{\mu} = \mu^*$ is learning perfectly. Then, we prove that $N \succeq \alpha I_2$ with $\alpha > 0$ for $\forall t \in [\delta, T]$ (see α in Eq.13).

Lemma 6.2. [Cross Term] Following setting of Lem. 6.1.

(a) **The** $|a_{ta} - v_{ta}| \leq \delta_{1,t}$ **situation.** For $\forall M(i, j)$, $|M(i, j)| \leq \gamma(\delta_{1,t})$, where $\gamma(\delta_{1,t}) \rightarrow 0$ as $\delta_{1,t} \rightarrow 0$ (see $\gamma(\delta_{1,t})$ in Eq.15).

(b) **The** $v_{ta} \geq a_{ta} + \delta_{1,t}$ **situation.** Let $\delta_{2,t} \triangleq v_{ta} - a_{ta} \geq \delta_{1,t}$ and $M_1 = M - M'$, where M' is SPD. Then, there exists an interval $v_{ta} \in [a_{ta} + \delta_{1,t}, a_{ta} + \delta_{2,t}]$ satisfies:

$$\begin{aligned} \mathbb{E}[M_1(1, 2)] &= \mathbb{E}[M_1(2, 1)] < 0, \\ \mathbb{E}[M_1(1, 1)] &= \mathbb{E}[M_1(2, 2)] > 0, \\ \mathbb{E}[M_1(1, 1) + M_1(1, 2)] &\geq u_1(v_{ta}, t) + u_2(v_{ta}, t), \end{aligned}$$

where $(u_1(v_{ta}, t) + u_2(v_{ta}, t))|_{v_{ta}=a_{ta}+\delta_{1,t}} > 0$, $u_1(\cdot, t)$ increasing and $u_2(\cdot, t)$ decreasing for $v_{ta} \in [a_{ta} + \delta_{1,t}, a_{ta} + \delta_{2,t}]$ (see M' , $u_1(\cdot, t)$ and $u_2(\cdot, t)$ in Eq. 16, 17 and 18).

To characterize the landscape of the objective function, we give the following definition.

Definition 6.3. $\phi : \mathbb{R}^D \rightarrow \mathbb{R}$ is λ -strongly convex and L_m -smooth if $\lambda I_D \preceq \nabla^2 \phi(x) \preceq L_m I_D$.

Since the Hessian matrix $H = 2(M + N)$, if $\alpha \geq \gamma$, we know that $\mathcal{L}$ is $2(\alpha - \gamma)$ -strongly convex for $|v_{ta} - a_{ta}| \leq \delta_{1,t}$. As shown in Lem. 6.2 (a), γ is related to the initialization area, and we can determine a suitable initialization parameter $\delta_{1,t}$ to guarantee $\alpha \geq \gamma$. For the setting $v_{ta} \geq a_{ta} + \delta_{2,t}$, we only require $u_1(v_{ta}, t) + u_2(v_{ta}, t) \geq 0$. The following condition shows our requirement for initialization, and the example shows that the initialization requirement is easy to satisfy.

Condition 1. $\delta_{1,t}$ satisfies $\alpha \geq \gamma(\delta_{1,t})$, and $\delta_{2,t}$ satisfies $u_1(a_{ta} + \delta_{2,t}) + u_2(a_{ta} + \delta_{2,t}) > 0$.

Example 6.4. Considering $A_s = [0.1, 0.1]^\top$, $A_{ta} = [0.12, 0.12]^\top$, and $\mu^* = 4$. With a $t = 2$, to satisfy **Condition 1**, we require $v_{ta}^{(0)} \in \{[0.1, 0.5] \cup [-0.5, -0.1]\}$, where 0.5 is far away from a_{ta} .

With a similar idea, we prove $\mathcal{L}$ is $2(\alpha + \gamma + \zeta)$ -smooth with $\zeta \geq 0$ (see ζ in Eq. 19). Let $V_{ta}^{(0)}$ be the initialization and $V_{ta}^{(k)}$ be the k -th iteration of GD algorithm. Then, we have the following results.

Theorem 6.5. Assume Theorem 3.1, 5.1, $\hat{\mu} = \mu^*$ and $\delta_{1,t}, \delta_{2,t}$ satisfy **Condition 1**. Considering score matching function $\mathcal{L}$. When $v_{ta}^{(0)} \in \{[a_{ta} - \delta_{1,t}, a_{ta} + \delta_{2,t}] \cup [-a_{ta} - \delta_{2,t}, -a_{ta} + \delta_{1,t}]\}$, using gradient descent with learning rate $\eta = 1/(2\alpha + \zeta)$, we have

$$\begin{aligned} \|V_{ta}^{(k)} V_{ta}^{(k)\top} - A_{ta} A_{ta}^\top\|_F &\leq \left(\frac{\kappa - 1}{\kappa + 1}\right)^k (2a_{ta} + \delta_{2,t}) |v_{ta}^{(0)} - a_{ta}|, \end{aligned}$$

where $\kappa = (\alpha + \gamma + \zeta)/(\alpha - \gamma)$.

This result is the first convergence guarantee for few-shot diffusion models and explains why few-shot models converge quickly to the ground-truth parameter.

Extension to K -mode GMM with General Dimension D . While the symmetric 2-mode case offers explicit constants, the geometric intuition generalizes to a K -mode GMM in D dimensions. Following the notation in Appendix F, we demonstrate that the score function's Jacobian retains a benign structure. This ensures the curvature remains positive definite, even within complex multi-modal distributions.

Lemma 6.6. The Jacobian matrix J of the K -mode GMM score function admits a rank-1 perturbed identity form: $J = \alpha I_D + \beta V_{ta} x^\top$. Consequently, the expected Squared Term is positive definite near the ground truth:

$$N \succeq \lambda_{sq} I_D, \quad \text{where } \lambda_{sq} > 0$$

is determined by the minimum eigenvalue of rank-1 perturbed structure (see λ_{sq} in Eq. (27) and Lem. F.1).

For the cross term, we derive a spectral norm bound that explicitly relates the perturbation strength to the parameter error and data complexity.

Lemma 6.7. [Perturbation Bound] The spectral norm of the Cross Term scales linearly with the parameter error $\|\Delta\|_F = \|V_{ta} - A_{ta}\|_F$. Specifically:

$$\|M_{\text{cross}}\|_2 \leq C_{\text{cross}} \|\Delta\|_F,$$

where C_{cross} (defined in Eq. 26) explicitly quantifies the data complexity, scaling with the diffusion noise level σ_t^{-2} and the higher-order moments of the latent GMM.

Combining the curvature provided by the Squared Term and the perturbation bound of the Cross Term, we establish the general strong convexity condition. This theorem generalizes the convergence guarantee to arbitrary dimensions and mode configurations.

Theorem 6.8. Assume the pretraining phase yields an initialization $V_{ta}^{(0)}$ such that $\|\Delta\|_{\text{init}} < \frac{\lambda_{sq}}{C_{\text{cross}}}$ holds. Then, the Hessian is strictly positive definite:

$$\lambda_{\min}(H) \geq 2\lambda_{sq} - 2C_{\text{cross}} \|\Delta\|_{\text{init}} > 0,$$

Table 2: $D = 8, d = 5$, 5-modal GMM latent.

v_{ta}	λ_1	λ_2	λ'_1	λ'_2
0.07	-2.8e-2	-2.7e-2	2.5e-4	1.5e-3
0.2	-7.1e-3	-6.9e-3	0.033	0.034
0.3	-2.6e-2	-2.1e-2	0.0738	0.076
0.5	-2.5e-2	-1.5e-2	0.206	0.211

Thus, the optimization objective is locally strongly convex, guaranteeing linear convergence to the ground truth A_{ta} .

This generalization confirms that the efficiency of few-shot adaptation is a robust property of diffusion models on multi-modal distributions, provided the pretraining establishes a sufficiently accurate latent structure to overcome the data complexity captured by C_{cross} .

Intuitively, for high-dimensional latent spaces ($d > 1$) with K -mode GMMs, the Hessian still decomposes into a Squared term (representing curvature) and a Cross term (representing residual error). Since the cross term is proportional to the residual error, the robust positive curvature of the squared term dominates the perturbation of the cross term when the ground truth is approached. This ensures that the local optimization landscape remains strongly convex, a geometric property empirically corroborated by the strictly positive eigenvalues observed in our high-dimensional simulations (Tables 2 and 4).

For the analysis of more general data, we also discuss extension to general bounded support latent (Section A.2) and multi low-dimensional linear subspace (Section A.3).

Simulation Experiments. We verify the optimization landscape by computing Hessian eigenvalues with μ_t^* and different v_{ta} for both phases and report eigenvalues. As shown in Table 2, pre-training objective yields negative eigenvalues ($\lambda_{1,2}$), confirming its non-convex nature. In contrast, few-shot phase exhibits strictly positive eigenvalues ($\lambda'_{1,2}$), which empirically supports the fast convergence guarantee in Thm. 6.5. Additional simulations confirm that these results generalize to high-dimensional multi-modal settings (Section B).

Remark 6.9 (Influence of t). Sec. 5 and 6 show that a great enough prior information $\hat{\mu}$ is important. When $t \rightarrow +\infty$, the information of $\hat{\mu}$ gradually disappears $\hat{\mu}_t \rightarrow 0$, which indicates the optimization process will become more difficult. Our convergence guarantee also reflects this intuition. When $t \rightarrow +\infty$, α (Eq. 13) and γ (Eq. 15) become 0, and the strongly convex parameter becomes smaller.

Remark 6.10 (Practical Operation). Sec. 6 supports other two operating principles. The cross-term error coefficient C_{cross} is directly proportional to the complexity of the data distribution (e.g., latent moments, modal dispersion). The more complex the concept, the smaller the allowable initial error radius for optimization, which relies on high-quality

pre-trained weights. Meanwhile, a great pre-trained model can provide a good optimized landscape during few-shot phase without requiring too many fine-tuning techniques.

6.1 DISCUSSION

Limited Source and Target Data. For the pretraining phase, we assume the latent parameter is perfectly learned $\hat{\mu} = \mu^*$ in Sec. 6. In this part, we discuss the setting $\|\hat{\mu} - \mu^*\| \leq \epsilon_{\text{pre}}$ (given a pretraining dataset with n_s datapoints, ϵ_{pre} has the order of $n_s^{-2/d}$ [Chen et al., 2023]). As shown in Lem 5.3 and Thm. 5.4, when $\|\hat{\mu} - \mu^*\| \leq \epsilon_{\text{pre}}$ is small enough, $\|\hat{V}_{ta} \hat{V}_{ta}^\top - A_{ta} A_{ta}^\top\|_F \leq \text{poly}(\epsilon_{\text{pre}})$. For the few-shot phase, Yang et al. [2024] show that there is an additional $1/\sqrt{n_{ta}}$ error with n_{ta} target data. Hence, there is an additional $\text{Poly}(n_s^{-2/d}) + n_{ta}^{-1/2}$ in Thm. 6.5 with a limited data.

Criteria of Great Pre-trained Model. The performance is determined by the scale of pretrained data, the model size, and the optimization step. In Section A.1, we discuss the balance between these terms from the theoretical perspective. From the empirical perspective, a great pre-trained model is usually an over-parametrized NN and is trained with a large-scale and diverse dataset. The optimization process should choose a suitable step that enjoys a converged loss and has the ability to generalize. We can use the standard quality metric, such as FID, IS, Clip Score, Pickscore, etc., to avoid underfitting. For overfitting, one can sample multiple images with the same prompts and observe their diversity and measure with the generalization score [Zhang et al., 2025a].

7 CONCLUSION

This work aims to explain why few-shot diffusion models can achieve great performance with a few optimization steps. As a start, we first evaluate each operation of few-shot diffusion models and show that a bad pre-trained model heavily influences the few-shot phase from the empirical and theoretical perspectives. After that, we show that a great pretrained model provides a warm-up for the few-shot phase and makes the landscape of few-shot diffusion models strongly convex. As a result, we prove that the few-shot model can fast converge to the ground-truth parameters by using a standard optimization algorithm (such as gradient descent). Combined with the gradient analysis for the pretraining phase, for the first time, we explain why few-shot models only need a few optimization steps compared with the pretraining phase to achieve great performance.

Future Work. While we extend our analysis to k -modal GMM latents (Theorem F.3), a gap to real-world complexity remains. Future research could explore general bounded support or multiple low-dimensional manifolds, as discussed in Section A.

ACKNOWLEDGEMENT

The corresponding author Shuai Li is supported by National Natural Science Foundation of China (92570111).

References

- Andreas Blattmann, Tim Dockhorn, Sumith Kulal, Daniel Mendelevitch, Maciej Kilian, Dominik Lorenz, Yam Levi, Zion English, Vikram Voleti, Adam Letts, et al. Stable video diffusion: Scaling latent video diffusion models to large datasets. *arXiv preprint arXiv:2311.15127*, 2023.
- Tony Bonnaire, Raphaël Urfin, Giulio Biroli, and Marc Mézard. Why diffusion models don’t memorize: The role of implicit dynamical regularization in training. *arXiv preprint arXiv:2505.17638*, 2025.
- Bradley CA Brown, Anthony L Caterini, Brendan Leigh Ross, Jesse C Cresswell, and Gabriel Loaiza-Ganem. Verifying the union of manifolds hypothesis for image data. In *ICLR*, 2023.
- Stefano Bruno, Ying Zhang, Dong-Young Lim, Ömer Deniz Akyildiz, and Sotirios Sabanis. On diffusion-based generative models and their error bounds: The log-concave case with full convergence estimates. *arXiv preprint arXiv:2311.13584*, 2023.
- Minshuo Chen, Kaixuan Huang, Tuo Zhao, and Mengdi Wang. Score approximation, estimation and distribution recovery of diffusion models on low-dimensional data. *arXiv preprint arXiv:2302.07194*, 2023.
- Sitan Chen, Vasilis Kontonis, and Kulin Shah. Learning general gaussian mixtures with efficient score matching. *arXiv preprint arXiv:2404.18893*, 2024.
- Ziheng Cheng, Tianyu Xie, Shiyue Zhang, and Cheng Zhang. Provable sample-efficient transfer learning conditional diffusion models via representation learning. *arXiv preprint arXiv:2502.04491*, 2025.
- Kurtland Chua, Qi Lei, and Jason D Lee. How fine-tuning allows for effective meta-learning. *Advances in Neural Information Processing Systems*, 34:8871–8884, 2021.
- Hugo Cui and Lenka Zdeborová. High-dimensional asymptotics of denoising autoencoders. *arXiv preprint arXiv:2305.11041*, 2023.
- Simon S Du, Wei Hu, Sham M Kakade, Jason D Lee, and Qi Lei. Few-shot learning via learning the representation, provably. *arXiv preprint arXiv:2002.09434*, 2020.
- Yingqing Guo, Hui Yuan, Yukang Yang, Minshuo Chen, and Mengdi Wang. Gradient guidance for diffusion models: An optimization perspective. *arXiv preprint arXiv:2404.14743*, 2024.
- Andi Han, Wei Huang, Yuan Cao, and Difan Zou. On the feature learning in diffusion models. *arXiv preprint arXiv:2412.01021*, 2024a.
- Yinbin Han, Meisam Razaviyayn, and Renyuan Xu. Neural network-based score estimation in diffusion models: Optimization and generalization. *arXiv preprint arXiv:2401.15604*, 2024b.
- Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. *Advances in Neural Information Processing Systems*, 35:26565–26577, 2022.
- Nupur Kumari, Bingliang Zhang, Richard Zhang, Eli Shechtman, and Jun-Yan Zhu. Multi-concept customization of text-to-image diffusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1931–1941, 2023.
- Puheng Li, Zhong Li, Huishuai Zhang, and Jiang Bian. On the generalization properties of diffusion models. *arXiv preprint arXiv:2311.01797*, 2023.
- Xiang Li, Yixiang Dai, and Qing Qu. Understanding generalizability of diffusion models requires rethinking the hidden gaussian structure. *Advances in neural information processing systems*, 37:57499–57538, 2024.
- Minghua Liu, Ruoxi Shi, Linghao Chen, Zhuoyang Zhang, Chao Xu, Xinyue Wei, Hansheng Chen, Chong Zeng, Jiayuan Gu, and Hao Su. One-2-3-45++: Fast single image to 3d objects with consistent multi-view generation and 3d diffusion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10072–10083, 2024.
- Zhiheng Liu, Ruili Feng, Kai Zhu, Yifei Zhang, Kecheng Zheng, Yu Liu, Deli Zhao, Jingren Zhou, and Yang Cao. Cones: Concept neurons in diffusion models for customized generation. *arXiv preprint arXiv:2303.05125*, 2023.
- Taehong Moon, Moonseok Choi, Gayoung Lee, Jung-Woo Ha, and Juho Lee. Fine-tuning diffusion models with limited data. In *NeurIPS 2022 Workshop on Score-Based Methods*, 2022.
- Phillip Pope, Chen Zhu, Ahmed Abdelkader, Micah Goldblum, and Tom Goldstein. The intrinsic dimension of images and its impact on learning. *arXiv preprint arXiv:2104.08894*, 2021.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695, 2022.
- Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22500–22510, 2023.
- Kulin Shah, Sitan Chen, and Adam Klivans. Learning mixtures of gaussians using the ddpm objective. *arXiv preprint arXiv:2307.01178*, 2023.
- Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. *arXiv preprint arXiv:2011.13456*, 2020.

- Pascal Vincent. A connection between score matching and denoising autoencoders. *Neural computation*, 23(7):1661–1674, 2011.
- Peng Wang, Huijie Zhang, Zekai Zhang, Siyi Chen, Yi Ma, and Qing Qu. Diffusion models learn low-dimensional distributions via subspace clustering. *arXiv preprint arXiv:2409.02426*, 2024.
- Chendong Xiang, Fan Bao, Chongxuan Li, Hang Su, and Jun Zhu. A closer look at parameter-efficient tuning in diffusion models. *arXiv preprint arXiv:2303.18181*, 2023.
- Ruofeng Yang, Bo Jiang, Cheng Chen, Ruinan Jin, Baoxiang Wang, and Shuai Li. Few-shot diffusion models escape the curse of dimensionality. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- Huijie Zhang, Jinfan Zhou, Yifu Lu, Minzhe Guo, Peng Wang, Liyue Shen, and Qing Qu. The emergence of reproducibility and generalizability in diffusion models. *arXiv preprint arXiv:2310.05264*, 2023.
- Huijie Zhang, Zijian Huang, Siyi Chen, Jinfan Zhou, Zekai Zhang, Peng Wang, and Qing Qu. Understanding generalization in diffusion models via probability flow distance. *arXiv preprint arXiv:2505.20123*, 2025a.
- Zekai Zhang, Xiao Li, Xiang Li, Lianghe Shi, Meng Wu, Molei Tao, and Qing Qu. Generalization of diffusion models arises with a balanced representation space. *arXiv preprint arXiv:2512.20963*, 2025b.

Evaluating the Role of Great Pre-trained Diffusion Models in Few-shot Phase: Warm-up and Acceleration (Supplementary Material)

Ruofeng Yang^{*1}

Yongcan Li^{*1}

Bo Jiang¹

Cheng Chen²

Shuai Li^{†1,3}

¹Shanghai Jiao Tong University, Shanghai, China

²East China Normal University, Shanghai, China,

³Shanghai Innovation Institute, Shanghai, China,

A MORE DISCUSSION ON GREAT PRE-TRAINED MODELS AND GENERAL LATENT

In this part, we first discuss the criteria of great pre-trained models and the influence of the pre-trained data, model size, and optimization steps.

A.1 DISCUSSION ON THE GREAT PRE-TRAINED MODELS

As the performance (generalization or memorization) of pre-trained models is determined by the scale of data, the model size, and optimization steps, we discuss the different combinations of these components from a theoretical perspective, which is helpful in determining the criteria of great pre-trained models. After that, we discuss how to extend to a general GMM with latent and bounded support. Finally, we intuitively discuss how to extend the few-shot diffusion models analysis to general multi low-dimensional manifolds.

The pretrained data. As discussed in Section 6.1, the limited pretrained data introduced the estimation error $\text{Poly}(n_s^{-2/d})$ for the pretraining phase. From the theoretical perspective, we require that the imperfect learning error for the pretraining phase is not the dominant term, instead of perfect learning. More specifically, we require $\text{Poly}(n_s^{-2/d})$ to be smaller than the optimization error in Thm. 6.5 and the limited target data error $n_{ta}^{-1/2}$.

The balance between pretrained data, model size, and optimization step. The above discussion builds on the size of pre-trained data and NN matches (as shown in Theorem 2 of Chen et al. [2023], the size of NN is determined by n_s) and ignores the optimization steps. However, there exists a mismatch between the pretrained data, model size, and optimization step in the application. Based on Li et al. [2024], we discuss the influence of the following mismatch cases for the pre-trained models:

Case 1: large train data and small NN size.

In this setting, the NN tends to learn the Gaussian structure of pretrained data (the empirical mean and covariance) instead of learning the multi-modal information of data. This setting belongs to the underfitting bad pre-trained models since it can not greatly learn the knowledge of the source data.

Case 2: Overparameterized NN with different optimization steps.

When an NN is overparameterized, with a large enough optimization step, the NN will memorize the training data, leading to an overfitted, bad pre-trained model.

With a small optimization step, Zhang et al. [2023] show that the NN still learn the Gaussian structure of training data instead of total source data knowledge, which belongs to the underfitting bad pretrained model.

A.2 EXTENSION TO GENERAL BOUNDED LATENT.

For the general latent distribution, if only focusing on the convergence guarantee, one possible way is to use the kernel-based method with a general wide 2-layer NN (the number of neurons $m = \Theta(n_s)$) [Li et al., 2023]:

$$s_{t,A}(X) := \text{AReLU}(WX + Ue(t))/m,$$

where $A \in \mathbb{R}^{D \times m}$ is trainable, $W \in \mathbb{R}^{m \times D}$ and $U \in \mathbb{R}^{m \times d_e}$ are randomly initialized and frozen during training, and $e(t)$ is the embedding of time. By setting $m = d$ (indicates d is large enough, which is also used by Han et al. [2024a]), the trainable A becomes the linear part, and the fixed $\text{ReLU}(WX + Ue(t))$ represents the nonlinear fixed latent in our work. Then, using the gradient flow algorithm, the score converges to the target distribution regardless of whether the pre-trained model is great (since the W, U are randomly initialized). Though this method can provide a convergence guarantee, it does not reflect the role of pretrained models and does not match the empirical operation. Hence, we adopt a simple setting to clearly explain the optimization process of the few-shot phase.

A.3 EXTENSION TO MULTI LOW-DIMENSIONAL SETTING

After obtaining the first convergence guarantee for the few-shot models under a single linear subspace with a GMM latent, we discuss how to extend the analysis to a union of low-dimensional subspaces.

Though real-world data admits the low-dimensional structure, it is a union of low-dimensional manifolds instead of one manifold [Brown et al., 2023]. Hence, a setting closer to the real-world data is to assume the target data is a union of linear subspaces. In the pretraining phase, Wang et al. [2024] makes the first step in this direction by modeling the data as a union of linear subspaces, and each subspace admits a Gaussian latent. We can first follow this direction and extend it to the few-shot phase. More specifically, for the few-shot modeling, we can assume the source and target data share some manifold and also have their own manifolds. Intuitively, since the pretraining phase has learn the shard manifold knowledge, based on our analysis, a great pre-trained model can also reduce the estimation error, warm-up, and accelerate the few-shot optimization process.

Go beyond: Few-shot analysis for a union of linear manifolds with general GMM latent. As Wang et al. [2024] assumes each manifold admits the Gaussian latent instead of the general GMM latent, it still has a gap to the real-world data. Another interesting future work is to combine the general GMM latent analysis (Theorem F.3) and multi-linear subspace assumption few-shot modeling to analyze the role of pre-trained models in the few-shot phase. We leave the analysis on the multi-subspace assumption and its GMM extension as interesting future works.

B ADDITIONAL EXPERIMENTS

B.1 ADDITIONAL SIMULATION EXPERIMENTS

In this part, we provide more simulation results with different D and d and show the two smallest eigenvalues of the Hessian matrix. As shown in the following two tables, the landscape of the pretraining phase is still non-convex. On the contrary, the landscape of the few-shot phase (with a great pretrained model) is almost strongly convex (except a very small negative eigenvalue $-7.5e - 5$).

The non-convex landscape of the pretraining phase indicates that it is possible to converge to the local minima instead of the global minima. We also verify this intuition through the simulation experiment. More specifically, we use the initialization area (v_{ta}, μ_t^*) ($v_{ta} = 0.07$ and ground truth $a_{ta} = 0.12$) and update models with GD algorithm. Then, the pretraining phase converges to the local minima 0.112, which is not equal to a_{ta} . On the contrary, the few-shot diffusion models with a fixed μ_t^* converge to 0.11999, almost the same as a_{ta} .

B.2 THE DETAIL OF THE UNDERFITTING REAL-WORLD EXPERIMENTS

In this part, we describe the setting of our real-world experiments on the CelebA 64 datasets. Our setting mainly follows Yang et al. [2024], and we provide the setting for the sake of completeness.

Value of v_{ta}	λ_1	λ_2	λ'_1	λ'_2
0.07	-0.0013	0.0015	0.0007	0.0016
0.2	-0.01	0.008	0.008	0.0083
0.3	-0.027	0.012	0.0126	0.0133
0.5	-0.057	0.013	0.0134	0.0151

Table 3: Eigenvalues for different v_{ta} ($D = 16, d = 1$)

Value of v_{ta}	λ_1	λ_2	λ'_1	λ'_2
0.07	-0.0002	-0.0002	-7.5e-5	1.24e-7
0.2	-0.0014	-0.0008	1.18e-5	1.29e-5
0.3	-0.0061	-0.0047	2.67e-5	2.86e-5
0.5	-0.0276	-0.0264	7.42e-5	7.89e-5

Table 4: Eigenvalues for different v_{ta} ($D = 4, d = 2$)

CelabA64 Datasets.

- Source dataset: 6400 images of faces with different hairstyles (without the bald feature).
- Target dataset: 10 images with the bald feature in CelebA64.

NN structure. In this experiment, we adopt a U-net network with 11 downblocks, 2 middleblocks, and 15 upblocks. In the pretraining phase, we train all parameters of the U-net. Since the NN layer in U-net is highly nonlinear, following Yang et al. [2024], we fine-tune the downblock and upblocks in the few-shot fine-tuning phase. More specifically, we fine-tune the first 4 downblock layers (as the encoder) and 4 upblock layers (as the decoder) in the fine-tuning phase.

The above experiments were conducted on a GeForce RTX 4090. For the pre-trained phase, we train the models for 50 epochs (bad pretrained model, Model-50) with batch size 20, which takes 1 hour. The great pretrained model (Model-200) takes 5 hours in the pretraining phase. For the fine-tuning phase, we fine-tune the pre-trained models with limited target datasets for 400 epochs with a batch size of 2. It takes 3 minutes to fine-tune the pre-trained models.

B.3 THE DETAIL OF THE OVERFITTING REAL-WORLD EXPERIMENTS

In the part, we provide the detail of the experiments on the Stable Diffusion models, including dataset and training pipeline.

Dataset and Evaluation Metric.

Training Dataset. The Dreambooth training dataset contains 30 subjects, and each subject contains 4-6 images to use to fine-tune the models (a total of 156 images).

Validation Dataset. The dreambooth dataset provides 25 test prompts for each subject (total $30 * 25 = 750$ prompts). Following the description of Dreambooth, we generate 4 images for each prompt and use these $3k$ images to evaluate.

Clip-T Score. Following Dreambooth, we calculate the cosine similarity between the prompt and image CLIP embeddings to measure the text-image alignment.

Pickscore. We also adopt the standard pickscore metric for text2image generation.

Training Dreambooth pipeline.

Overfitting Bad Models. Since the SD 1.4 and SD 3 Medium can generate high-quality and diverse samples, we view them as great pretrained models. To obtain a bad pretrained model, we overfit the SD 1.4 and SD 3 Medium with one prompt (a photo of a dog) and the corresponding 5 images. As the overfitting step increases, the diversity of pretrained models decreases (preferring to generate dog images in our setting). We use two bad pretrained models, the first one overfits the one prompt 1k steps, and the latter, a worse one, overfits 4k steps (lower diversity). The overfitting learning rate is 5×10^{-6} , the resolution is 512 for SD 1.4 and 1024 for SD3-Medium and the accumulation steps is 4.

Fine-tuning phase with freezing most parameters. Then, we modify the train Dreambooth pipeline of the diffuser to train with the training dataset. To match the setting of our theoretical results, we only fine-tune the first and last 3 blocks of NN (Unet of SD 1.4 and DiT of SD 3). The fine-tuning optimization step is $1k$. The learning rate and resolution is the same with the overfitting phase.

C THE DETAILED CALCULATION OF GRADIENT AND HESSIAN

Since our analysis depends heavily on the gradient and Hessian for the few-shot score matching objective function, we provide the detailed form of these terms in this section.

C.1 TERMS RELATED TO $\mathcal{L}_{SM,t}^{\text{few}}$

Recall that $\mathbb{E} \frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2}$ is consisted by the cross and squared term

$$\mathbb{E} \frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2} = 2 \left[\mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] + \mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right] \right].$$

To obtain the expectation form of $\mathbb{E} \frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2}$, we calculate the expectation form of each term.

Calculate $\frac{\partial \mathcal{L}_{SM,t}^{\text{few}}}{\partial V_{ta}}$. For this term, we know that

$$\frac{\partial \mathcal{L}_{SM,t}^{\text{few}}}{\partial V_{ta}} = 2 \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top (s_{\hat{\mu}, V_{ta}}(x_t, t) - s_{\mu^*, A_{ta}}(x_t, t)),$$

where

$$s_{\mu^*, A}(x_t, t) = A \tanh(\mu_t^*{}^\top A^\top x_t) \mu_t^* - A A^\top x_t - \frac{1}{\sigma_t^2} (I_D - A A^\top) x_t.$$

For the first term, we have the following equation:

$$\begin{aligned} \frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} &= \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t I_2 + \frac{\partial \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t}{\partial V_{ta}} V_{ta}^\top + \left(\frac{1}{\sigma_t^2} - 1 \right) \frac{\partial V_{ta} V_{ta}^\top x_t}{\partial V_{ta}} \\ &= \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t I_2 + (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t V_{ta}^\top + \left(\frac{1}{\sigma_t^2} - 1 \right) (x_t V_{ta}^\top + V_{ta}^\top x_t I_2) \\ &= \left(\tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1 \right) V_{ta}^\top x_t \right) I_2 + \left((1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1 \right) \right) x_t V_{ta}^\top. \end{aligned}$$

Calculate $\frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2}$. We know that the Hessian matrix of the few-shot score matching objective function can be decomposed into the cross term and the squared term.

$$\frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2} = 2 \underbrace{\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)}_{\text{Squared Term } N} + 2 \underbrace{\left(\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}})}_{\text{Cross Term } M}.$$

For the cross term, we know that

$$\begin{aligned}
& \frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \\
&= \begin{bmatrix} (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t \begin{bmatrix} x_t(1) & 0 \\ 0 & x_t(1) \end{bmatrix} \\ (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t \begin{bmatrix} x_t(2) & 0 \\ 0 & x_t(2) \end{bmatrix} \end{bmatrix} + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) \begin{bmatrix} \begin{bmatrix} x_t(1) & 0 \\ 0 & x_t(1) \end{bmatrix} \\ \begin{bmatrix} x_t(2) & 0 \\ 0 & x_t(2) \end{bmatrix} \end{bmatrix} \\
&+ \begin{bmatrix} (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t \begin{bmatrix} x_t(1) & 0 \\ x_t(2) & 0 \end{bmatrix} \\ (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t \begin{bmatrix} 0 & x_t(1) \\ 0 & x_t(2) \end{bmatrix} \end{bmatrix} \\
&- 2 \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t \begin{bmatrix} \hat{\mu}_t x_t(1) x_t V_{ta}^\top \\ \hat{\mu}_t x_t(2) x_t V_{ta}^\top \end{bmatrix}.
\end{aligned}$$

Let $s_{\hat{\mu}, V_{ta}}(x_t, t) - s_{\mu^*, A_{ta}}(x_t, t) = y$. Then, we have that

$$\begin{aligned}
& \left(\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \\
&= (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y I_2 + (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t y^\top \\
&- 2 \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y x_t V_{ta}^\top + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) x_t^\top y I_2.
\end{aligned}$$

For the squared term, we know that

$$\begin{aligned}
& \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) = \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t I_2 + (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t))^2 \hat{\mu}_t^\top \hat{\mu}_t V_{ta} x_t^\top x_t V_{ta}^\top \\
&+ \left(\frac{1}{\sigma_t^2} - 1 \right)^2 (V_{ta} x_t^\top x_t V_{ta}^\top + x_t^\top V_{ta} V_{ta}^\top x_t I_2 + V_{ta} x_t^\top V_{ta}^\top x_t + x_t^\top V_{ta} x_t V_{ta}^\top) \\
&+ 2(1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t V_{ta} x_t^\top \\
&+ 2 \left(\frac{1}{\sigma_t^2} - 1 \right) \tanh(\mu^\top V_{ta}^\top x_t) \hat{\mu}_t (x_t V_{ta}^\top + V_{ta}^\top x_t I_2) \\
&+ 2(1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \left(\frac{1}{\sigma_t^2} - 1 \right) \hat{\mu}_t^\top \hat{\mu}_t V_{ta}^\top V_{ta} x_t x_t^\top \\
&+ (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \left(\frac{1}{\sigma_t^2} - 1 \right) \hat{\mu}_t^\top \hat{\mu}_t (x_t x_t^\top V_{ta} V_{ta}^\top + V_{ta} V_{ta}^\top x_t x_t^\top).
\end{aligned}$$

Calculate the expectation of Hessian $\mathbb{E} \frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2}$. As discussed in Section 6.1, we take expectation over the target distribution q_{ta} . Hence, we calculate the expectation of Hessian.

Before providing the result of the Hessian matrix, we first do some helpful calculation. Recall that under the linear subspace assumption (Theorem 3.1), the diffusion process happens in the latent distribution $z_0 \sim \frac{1}{2} \mathcal{N}(\mu^*, 1) + \frac{1}{2} \mathcal{N}(-\mu^*, 1)$, which indicates $z_t = \exp(-t)z_0 + \sqrt{1 - \exp(-2t)}\xi_t$ with $\xi_t \sim \mathcal{N}(0, 1)$. Then, by changing the probability density variable, we have

$$\begin{aligned}
& \exp(-t)z_0 = \frac{1}{2} \mathcal{N}(\exp(-t)\mu^*, \exp(-2t)) + \frac{1}{2} \mathcal{N}(-\exp(-t)\mu^*, \exp(-2t)) \\
& \sqrt{1 - \exp(-2t)}\xi_t \sim \mathcal{N}(0, (1 - \exp(-2t))) \\
& z_t = \exp(-t)z_0 + \sqrt{1 - \exp(-2t)}\xi_t \sim \frac{1}{2} \mathcal{N}(\mu_t^*, 1) + \frac{1}{2} \mathcal{N}(-\mu_t^*, 1).
\end{aligned}$$

Then, we know that $z_t \sim N(\mu_t^*, 1)$, where $\mu_t^* = \exp(-t)\mu^*$, which indicates

$$x_t = A_{ta} z_t \sim \frac{1}{2} N(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top) + \frac{1}{2} N(-\mu_t^* A_{ta}, A_{ta} A_{ta}^\top),$$

and

$$V_{ta}^\top x_t \sim \frac{1}{2}N(\mu_t^* V_{ta}^\top A_{ta}, V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) + \frac{1}{2}N(-\mu_t^* V_{ta}^\top A_{ta}, V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}).$$

We should first calculate $\mathbb{E}[x_t x_t^\top]$, $\mathbb{E}[x_t^\top x_t]$, $\mathbb{E}[x_t^\top y]$ and $\mathbb{E}[x_t y^\top]$, where $y = s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}$, as they will be frequently utilized in subsequent steps.

$$\begin{aligned} \mathbb{E}[x_t^\top x_t] &= \mathbb{E}[\Sigma x_t(i)^2] = \Sigma D[x_t(i)] + \mathbb{E}^2[x_t(i)] \\ &= \text{tr}(A_{ta} A_{ta}^\top) + \mathbb{E}[x_t]^\top \mathbb{E}[x_t] \\ &= \text{tr}(A_{ta}^\top A_{ta}) + \mathbb{E}[x_t]^\top \mathbb{E}[x_t] \\ &= (1 + \mu_t^{*2}) A_{ta}^\top A_{ta}. \end{aligned} \tag{4}$$

$$\begin{aligned} \mathbb{E}[x_t x_t^\top] &= \mathbb{E}[(x_t - \mu^* A_{ta})(x_t - \mu^* A_{ta})^\top] + \mu^* A_{ta}^* \mathbb{E}[x_t^\top] + \mu^* \mathbb{E}[x_t] A_{ta}^\top - \mu^{*2} A_{ta} A_{ta}^\top \\ &= (1 + \mu^{*2}) A_{ta} A_{ta}^\top \end{aligned}$$

Observe that x_t is a symmetric distribution. Then, for any even function f , we can write

$$\begin{aligned} \mathbb{E}_{x_t}[f(x_t)] &= \frac{1}{2} \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[f(x_t)] + \frac{1}{2} \mathbb{E}_{x_t \sim \mathcal{N}(-\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[f(x_t)] \\ &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[f(x_t)]. \end{aligned}$$

Applying this property of the even function, we can obtain the following result by using the fact that $x_t^\top y$ and $x_t y^\top$ are even functions x_t (recall that $s_{\hat{\mu}, V_{ta}}(x_t, t) - s_{\mu^*, A_{ta}}(x_t, t) = y$).

$$\begin{aligned} \mathbb{E}[x_t^\top y] &= \mathbb{E}_{x_t}[x_t^\top (V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t - V_{ta} V_{ta}^\top x_t - \frac{1}{\sigma_t^2} (I_2 - V_{ta} V_{ta}^\top) x_t)] - \mathbb{E}_{x_t}[x_t^\top s_{\mu^*, A_{ta}}] \\ &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t + (\frac{1}{\sigma_t^2} - 1) x_t^\top V_{ta} V_{ta}^\top x_t - \frac{1}{\sigma_t^2} x_t^\top x_t] \\ &\quad - \mathbb{E}_{x_t}[x_t^\top s_{\mu^*, A_{ta}}] \end{aligned} \tag{5}$$

$$\begin{aligned} &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t] + (\frac{1}{\sigma_t^2} - 1) ((1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) \\ &\quad - \frac{1}{\sigma_t^2} (\text{tr}(A_{ta} A_{ta}^\top) + \hat{\mu}_t^2 A_{ta}^\top A_{ta}) - \mathbb{E}_{x_t}[x_t^\top s_{\mu^*, A_{ta}}] \\ &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t] + (\frac{1}{\sigma_t^2} - 1) ((1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) \\ &\quad - \frac{1}{\sigma_t^2} (\text{tr}(A_{ta}^\top A_{ta}) + \hat{\mu}_t^2 A_{ta}^\top A_{ta}) - \mathbb{E}_{x_t}[x_t^\top s_{\mu^*, A_{ta}}] \\ &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t] + (\frac{1}{\sigma_t^2} - 1) ((1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) \\ &\quad - \frac{1}{\sigma_t^2} (\text{tr}(A_{ta}^\top A_{ta}) + \hat{\mu}_t^2 A_{ta}^\top A_{ta}) - \mathbb{E}_{x_t}[x_t^\top s_{\mu^*, A_{ta}}] \\ &= \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t] + (\frac{1}{\sigma_t^2} - 1) ((1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) \\ &\quad - \mathbb{E}_{x_t \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)}[x_t^\top A_{ta} \tanh(\mu_t^{* \top} A_{ta}^\top x_t) \mu_t^*] - (\frac{1}{\sigma_t^2} - 1) ((1 + \mu_t^{*2}) A_{ta}^\top A_{ta} A_{ta}^\top A_{ta}). \end{aligned} \tag{6}$$

Through similar calculations, we can also get $\mathbb{E}_{x_t}[x_t y^\top]$:

$$\begin{aligned}
\mathbb{E}_{x_t}[x_t y^\top] &= \mathbb{E}[x_t (V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t - x_t^\top V_{ta} V_{ta}^\top - \frac{1}{\sigma_t^2} x_t^\top (I_2 - V_{ta} V_{ta}^\top) \\
&\quad - A_{ta} \tanh(\hat{\mu}_t^\top A_{ta}^\top x_t) \hat{\mu}_t - x_t^\top A_{ta} A_{ta}^\top - \frac{1}{\sigma_t^2} x_t^\top (I_2 - A_{ta} A_{ta}^\top))] \\
&= \mathbb{E}_{x_t \sim N(\mu_t^*, A_{ta}, A_{ta}^\top)} [x_t (V_{ta} \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t - x_t (A_{ta} \tanh(\mu_t^{*\top} A_{ta}^\top x_t) \mu_t^*) \\
&\quad + (1 + \mu_t^{*2}) (\frac{1}{\sigma_t^2} - 1) A_{ta} A_{ta}^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top)]. \tag{7}
\end{aligned}$$

With the calculation for $\mathbb{E}[x_t^\top y]$ and $\mathbb{E}_{x_t}[x_t y^\top]$, we can obtain the expectation form of the cross and squared term. For the cross term, we have that

$$\begin{aligned}
&\mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] \\
&= \mathbb{E} \left[(1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y I_2 + (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t y^\top \right. \\
&\quad \left. - 2 \tanh(\mu_t^\top V_{ta}^\top x_t) (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y x_t V_{ta}^\top + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) x_t^\top y I_2 \right] \\
&= \mathbb{E} [\hat{\mu}_t^\top \hat{\mu}_t x_t^\top y I_2] - \mathbb{E} [\tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y] + \mathbb{E} [\hat{\mu}_t^\top \hat{\mu}_t x_t y^\top] \\
&\quad - \mathbb{E} [\tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t x_t y^\top] \\
&\quad - 2 \mathbb{E} [\tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y x_t V_{ta}^\top] + 2 \mathbb{E} [\tanh^3(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t x_t^\top y x_t V_{ta}^\top] \\
&\quad + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) \mathbb{E} [x_t^\top y I_2].
\end{aligned}$$

For the squared term, we have that

$$\begin{aligned}
&\mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right] \\
&= \mathbb{E} [\tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t I_2] + \mathbb{E} [(1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t))^2 \hat{\mu}_t^\top \hat{\mu}_t V_{ta} x_t^\top x_t V_{ta}^\top] \\
&\quad + 2 \left(\frac{1}{\sigma_t^2} - 1 \right)^2 ((1 + \hat{\mu}_t^2) A_{ta} A_{ta}^\top V_{ta} V_{ta}^\top + (1 + \hat{\mu}_t^2) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta} I_2) \\
&\quad + \mathbb{E} [2(1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t^\top \hat{\mu}_t V_{ta} x_t^\top] \\
&\quad + 2 \mathbb{E} \left[\left(\frac{1}{\sigma_t^2} - 1 \right) \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t (x_t V_{ta}^\top + V_{ta}^\top x_t I_2) \right] \\
&\quad + \mathbb{E} (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \left(\frac{1}{\sigma_t^2} - 1 \right) \hat{\mu}_t^\top \hat{\mu}_t (2 V_{ta}^\top V_{ta} x_t x_t^\top + x_t x_t^\top V_{ta} V_{ta}^\top + V_{ta} V_{ta}^\top x_t x_t^\top).
\end{aligned}$$

C.2 TERMS RELATED TO $\mathcal{L}_{\text{SM}}^{\text{few-all}}$

For the fully fine-tuning method, we show that $\frac{\partial \mathcal{L}_{\text{SM}}^{\text{few-all}}}{\partial \mu_t}$ is small in Theorem 5.5. In this part, we provide the calculation of this term. We note that when considering fully fine-tuning method, μ_t also has a gradient.

$$\frac{\partial \mathcal{L}_{\text{SM}}^{\text{few-all}}}{\partial \mu_t} = 2 (s_{\mu, V_{ta}}(x_t, t) - s_{\mu^*, A_{ta}})^\top (V_{ta} \tanh(\mu_t^\top V_{ta}^\top x_t) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x_t)) V_{ta}^\top x_t).$$

D THE PROOF FOR BAD PRETRAINING

Lemma D.1. Assume Theorem 3.1 and 5.1. If $\hat{\mu} \neq \mu^*$, with $V_{ta} V_{ta}^\top = A_{ta} A_{ta}^\top$, $\partial \mathcal{L} / \partial V_{ta} \neq 0$.

Proof. For the sake of brevity, we use x, μ_t instead of $x_t^{ta}, \hat{\mu}_t$ when there is no ambiguity in this part. We also ignore (x_t, t) in $s_{\hat{\mu}, V_{ta}}(x_t, t)$ and $s_{\mu^*, A_{ta}}(x_t, t)$ for clarity.

We know that

$$\mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}} \right] = \mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top (s_{\hat{\mu}, V_{ta}}(x, t) - s_{\mu^*, A_{ta}}(x, t)) \right].$$

For each term, we have the following form:

$$\begin{aligned} \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top &= \tanh(\mu_t V_{ta}^\top x) \mu_t I_2 + (1 - \tanh^2(\mu_t V_{ta}^\top x)) \mu_t^\top \mu_t V_{ta} x^\top \\ &\quad + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2) \\ &\triangleq f(x, V_{ta}, \mu_t) + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2), \end{aligned}$$

and

$$\begin{aligned} s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}} &= V_{ta} \mu_t \tanh(\mu_t V_{ta}^\top x) - A_{ta} \mu_t^* \tanh(\mu_t^* A_{ta}^\top x) + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \\ &\triangleq g(x, V_{ta}, \mu_t, \mu_t^*) + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x. \end{aligned}$$

Then, we simplify the gradient term into the following form

$$\begin{aligned} \mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}} \right] &= \mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] \\ &= \mathbb{E} \left[\left(f(x, V_{ta}, \mu_t) + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2) \right) g(x, V_{ta}, \mu_t, \mu_t^*) \right] \\ &\quad + \mathbb{E} \left[\left(f(x, V_{ta}, \mu_t) + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2) \right) \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \right] \\ &= \mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right)^2 (V_{ta} x^\top + V_{ta}^\top x I_2) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \right] + \mathbb{E}_x [h(x, V_{ta}, A_{ta}, \mu_t, \mu_t^*)], \end{aligned}$$

where

$$\begin{aligned} h(x, V_{ta}, A_{ta}, \mu_t, \mu_t^*) &= f \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2) g + f g \\ &= \left(\frac{1}{\sigma_t^2} - 1 \right) (\tanh(\mu_t V_{ta}^\top x) \mu_t I_2 + (1 - \tanh^2(\mu_t V_{ta}^\top x)) \mu_t^2 V_{ta} x^\top) (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \\ &\quad + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta} x^\top + V_{ta}^\top x I_2) (V_{ta} \mu_t \tanh(\mu_t V_{ta}^\top x) - A_{ta} \mu_t^* \tanh(\mu_t^* A_{ta}^\top x)) \\ &\quad + (\tanh(\mu_t V_{ta}^\top x) \mu_t I_2 \\ &\quad + (1 - \tanh^2(\mu_t V_{ta}^\top x)) \mu_t^2 V_{ta} x^\top) (V_{ta} \mu_t \tanh(\mu_t V_{ta}^\top x) - A_{ta} \mu_t^* \tanh(\mu_t^* A_{ta}^\top x)). \end{aligned} \tag{8}$$

We first calculate $\mathbb{E}_x [V_{ta} x^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x]$ and $\mathbb{E}_x [(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x x^\top V_{ta}]$, which is useful in bounding the first term of $\mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}} \right]$:

$$\begin{aligned} \mathbb{E}_x [V_{ta} x^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x] &= V_{ta} \mathbb{E}_x [x^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x] \\ &= V_{ta} \mathbb{E}_x [\text{tr}(x^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x)] \\ &= V_{ta} \mathbb{E}_x [\text{tr}((V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x x^\top)] \\ &= V_{ta} \text{tr}(\mathbb{E}_x [(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x x^\top]) \\ &= (1 + \mu_t^{*2}) \text{tr}((V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) A_{ta} A_{ta}^\top) V_{ta}, \end{aligned}$$

where the last equality follows the fact that $\mathbb{E}[xx^\top] = (1 + \mu_t^{*2})A_{ta}A_{ta}^\top$ (Eq.4). Similarly, we can obtain the following bound:

$$\begin{aligned}\mathbb{E}_x [(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)xx^\top V_{ta}] &= (V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)\mathbb{E}_X[xx^\top]V_{ta} \\ &= (1 + \mu_t^{*2})(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top V_{ta}.\end{aligned}$$

Thus, the first term of the gradient $\mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}} \right]$ has the following form:

$$\begin{aligned}&\mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right)^2 (V_{ta}x^\top + V_{ta}^\top x I_2)(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)x \right] \\ &= \left(\frac{1}{\sigma_t^2} - 1 \right)^2 \mathbb{E}_x [(V_{ta}x^\top + V_{ta}^\top x I_2)(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)x] \\ &= \left(\frac{1}{\sigma_t^2} - 1 \right)^2 (\mathbb{E}_x [V_{ta}x^\top (V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)x] + \mathbb{E}_x [(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)xx^\top V_{ta}]) \\ &= \left(\frac{1}{\sigma_t^2} - 1 \right)^2 (1 + \mu_t^{*2})((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top V_{ta} + \text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top)I_2)V_{ta} \\ &\triangleq w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top)V_{ta}.\end{aligned}\tag{9}$$

Let $-\mathbb{E}_x[h(x, V_{ta}, A_{ta}, \mu_t, \mu_t^*)] \triangleq h(V_{ta}, A_{ta}, \mu_t, \mu_t^*)$, we know that

$$\mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] = 0$$

is equivalent to

$$w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top)V_{ta} = h(V_{ta}, A_{ta}, \mu_t, \mu_t^*).$$

We then prove that $w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top) = 0$ if and only if $V_{ta}V_{ta}^\top = A_{ta}A_{ta}^\top$ when $A_{ta} \neq 0$.

If $V_{ta}V_{ta}^\top = A_{ta}A_{ta}^\top$:

$$\begin{aligned}V_{ta}V_{ta}^\top &= A_{ta}A_{ta}^\top \Rightarrow V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top = 0 \\ &\Rightarrow w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top) = 0\end{aligned}$$

If $w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top) = 0$, we know that

$$\begin{aligned}(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top + \text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top)I_2 &= 0 \\ \Rightarrow \text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top) &= -2\text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top),\end{aligned}$$

which indicates $\text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top) = 0$. Then, we know that

$$\begin{aligned}V_{ta}V_{ta}^\top A_{ta}A_{ta}^\top - A_{ta}A_{ta}^\top A_{ta}A_{ta}^\top &= -\text{tr}((AA^\top - A_{ta}A_{ta}^\top)A_{ta}A_{ta}^\top)I_2 = 0 \\ \Rightarrow \text{tr}(V_{ta}V_{ta}^\top A_{ta}A_{ta}^\top) &= V_{ta}^\top A_{ta}A_{ta}^\top V_{ta} = A_{ta}^\top A_{ta}A_{ta}^\top A_{ta} \\ \Rightarrow V_{ta}^\top A_{ta} &= \pm A_{ta}^\top A_{ta} \\ \Rightarrow V_{ta}A_{ta}^\top A_{ta}A_{ta}^\top &= \pm A_{ta}A_{ta}^\top A_{ta}A_{ta}^\top \\ \Rightarrow V_{ta}A_{ta}^\top &= \pm A_{ta}A_{ta}^\top \\ \Rightarrow V_{ta} &= \pm A_{ta}, V_{ta}V_{ta}^\top = A_{ta}A_{ta}^\top\end{aligned}$$

Then we need $h(V_{ta}, A_{ta}, \mu_t, \mu_t^*) = 0$. However, if $\mu_t \neq \mu_t^*$, $h(V_{ta}, A_{ta}, \mu_t, \mu_t^*) \neq 0$ when $V_{ta}V_{ta}^\top = A_{ta}A_{ta}^\top$. In other words, if $\mu_t \neq \mu_t^*$, $V_{ta}V_{ta}^\top = A_{ta}A_{ta}^\top$ can not make $\mathbb{E} \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}} \right] = 0$. Then, we complete the proof. $\blacksquare$

Theorem 5.4. Assume Theorem 3.1 and 5.1 hold. Let μ_1^* and μ_2^* be the two parameters to generate different latent distributions. Given a bad pre-trained model with $\hat{\mu} = 0$, if $|\mu_1^* - \hat{\mu}| > |\mu_2^* - \hat{\mu}|$, then

$$\|\widehat{V}_{ta,1}\widehat{V}_{ta,1}^\top - A_{ta}A_{ta}^\top\|_F > \|\widehat{V}_{ta,2}\widehat{V}_{ta,2}^\top - A_{ta}A_{ta}^\top\|_F,$$

where $\widehat{V}_{ta,i}$ is the solution corresponds to μ_i^* , $i \in \{1, 2\}$.

Proof. With $\mu_t = 0$ and $\mu_t^* \neq 0$, then we know that

$$h(V_{ta}, A_{ta}, \mu_t, \mu_t^*) = \mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta}x^\top + V_{ta}^\top x I_2) \tanh(\mu_t^* A_{ta}^\top x) \mu_t^* \right] A_{ta}.$$

We know that

$$\mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right) (x^\top A_{ta} + A_{ta}x^\top) \tanh(\mu_t^* A_{ta}^\top x) \mu_t^* \right] > 0$$

and

$$\mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right) (x^\top A_{ta} + A_{ta}x^\top) \tanh(\mu_t^* A_{ta}^\top x) \mu_t^* \right] > 0,$$

so $w(V_{ta}V_{ta}^\top, A_{ta}A_{ta}^\top) > 0$, which means that there exists a constant positive gap between $V_{ta}V_{ta}^\top$ and $A_{ta}A_{ta}^\top$.

We also know that function $x \tanh x$ is even, which indicates if $\mu_1^* > \mu_2^*$,

$$\begin{aligned} & \mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right) (x^\top A_{ta} + A_{ta}x^\top) \tanh(\mu_1^* A_{ta}^\top x) \mu_1^* \right] \\ & > \mathbb{E}_x \left[\left(\frac{1}{\sigma_t^2} - 1 \right) (x^\top A_{ta} + A_{ta}x^\top) \tanh(\mu_2^* A_{ta}^\top x) \mu_2^* \right]. \end{aligned}$$

Therefore, the constant positive gap between $V_{ta}V_{ta}^\top$ and $A_{ta}A_{ta}^\top$ must increase.

$$\|\widehat{V}_{ta,1}\widehat{V}_{ta,1}^\top - A_{ta}A_{ta}^\top\|_F > \|\widehat{V}_{ta,2}\widehat{V}_{ta,2}^\top - A_{ta}A_{ta}^\top\|_F$$

■

Theorem 5.5. Assume Theorem 3.1 and 5.1 holds. For a fixed t , if $\mu_t \in (-\epsilon, \epsilon)$, we have that

$$\begin{aligned} \partial \mathcal{L}_{\text{SM},t}^{\text{few-all}} / \partial \mu_t & \leq \\ & 4\epsilon A_{ta}^\top V_{ta} \sqrt{(1 + \mu_t^{*2}) V_{ta}^\top V_{ta} \sqrt{C_1} + O(\epsilon^{\frac{3}{2}})}, \end{aligned}$$

where C_1 is a small constant determined by V_{ta} , A_{ta} and μ^* (Details in Eq. 10).

Proof. Through simple algebraic calculations, we know the gradient for μ_t have the following form:

$$\begin{aligned} \frac{\partial \mathcal{L}_{\text{SM}}^{\text{few-all}}}{\partial \mu_t} & = 2(\hat{s}_{\mu, V_{ta}} - s_{\mu^*, A_{ta}})^\top (V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x) \\ & = 2y^\top (V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x). \end{aligned}$$

For the term, by using the Cauchy-Schwarz inequality, we know that

$$\begin{aligned} & \mathbb{E}_x [2y^\top (V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x)] \\ & = \mathbb{E}_{x \sim \mathcal{N}(\mu_t^* A_{ta}, A_{ta} A_{ta}^\top)} [2y^\top (V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x)] \\ & \leq 2\sqrt{\mathbb{E}[y^\top y]} \times \\ & \sqrt{\mathbb{E}[\|V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x\|_2^2]}. \end{aligned}$$

Then we give the upper bounds on $\mathbb{E}[y^\top y]$ and

$$\mathbb{E} \left[\|V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x\|_2^2 \right]$$

to achieve the final bound.

For the second part, if $\mu_t \in (-\epsilon, \epsilon)$, we have

$$\begin{aligned} & \left[\|V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) + \mu_t V_{ta} (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) V_{ta}^\top x\|_2 \right] \\ & \leq \mathbb{E} [\epsilon^2 V_{ta}^\top V_{ta} V_{ta}^\top x x^\top V_{ta} + \epsilon^2 x^\top V_{ta} V_{ta}^\top V_{ta} V_{ta}^\top x + 2\epsilon^2 x^\top V_{ta} V_{ta}^\top V_{ta} V_{ta}^\top x] \\ & = 4\epsilon^2 (1 + \mu_t^{*2}) V_{ta}^\top V_{ta} V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}, \end{aligned}$$

where the inequality follows by the fact $\tanh^2(\mu_t^\top V_{ta}^\top x) \in [0, 1]$ and the first order of Taylor expansion for $\tanh(\mu_t^\top V_{ta}^\top x)$ (when $\mu_t \in (-\epsilon, \epsilon)$ is close to 0, the influence of higher-order terms in Taylor expansion is limited). The last equality follows Equation (4).

For the first part, we can divide $\mathbb{E}[y^\top y]$ into three parts below:

$$\begin{aligned} \mathbb{E}[y^\top y] &= \mathbb{E} [\|V_{ta} \tanh(\mu_t V_{ta}^\top x) - A_{ta} \tanh(\mu_t^* A_{ta}^\top x)\|_2^2] \\ &+ 2\mathbb{E}[(\tanh(\mu_t V_{ta}^\top x) V_{ta}^\top - \tanh(\mu_t^* A_{ta}^\top x) A_{ta}^\top)(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x] \\ &+ \left(\frac{1}{\sigma_t^2} - 1 \right)^2 \text{tr}((V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top)^2 A_{ta} A_{ta}^\top). \end{aligned}$$

Next we bound each of these three terms separately.

Bound for $\mathbb{E}[\|V_{ta} \tanh(\mu_t V_{ta}^\top x) - A_{ta} \tanh(\mu_t^* A_{ta}^\top x)\|_2^2]$.

$$\begin{aligned} & \mathbb{E} [\|V_{ta} \tanh(\mu_t V_{ta}^\top x) - A_{ta} \tanh(\mu_t^* A_{ta}^\top x)\|_2^2] \\ & \leq \mathbb{E} [\epsilon^2 V_{ta}^\top V_{ta} V_{ta}^\top x x^\top V_{ta} + \mu_t^{*2} A_{ta}^\top A_{ta} A_{ta}^\top x x^\top A_{ta} + 2\epsilon \mu_t^* V_{ta}^\top A_{ta} V_{ta}^\top x x^\top A_{ta}] \\ & = \left((\epsilon^2 V_{ta}^\top V_{ta} V_{ta}^\top A_{ta} A_{ta}^\top V_{ta} + \mu_t^{*2} A_{ta}^\top A_{ta} A_{ta}^\top A_{ta} A_{ta}^\top A_{ta} + 2\epsilon \mu_t^* V_{ta}^\top A_{ta} A_{ta}^\top A_{ta} A_{ta}^\top A_{ta}) \right. \\ & \quad \left. \times (1 + \mu_t^{*2}) \right) \triangleq C_1, \end{aligned}$$

where the inequality follows by Equation (4) and the first order of Taylor expansion for $\tanh(\mu_t^\top V_{ta}^\top x)$ (when $\mu_t \in (-\epsilon, \epsilon)$ is close to 0, the influence of higher-order terms in Taylor expansion is limited).

Bound for $2\mathbb{E}[(\tanh(\mu_t V_{ta}^\top x) V_{ta}^\top - \tanh(\mu_t^* A_{ta}^\top x) A_{ta}^\top)(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x]$.

By simple algebraic calculation, we know that

$$\begin{aligned} & 2(\tanh(\mu_t V_{ta}^\top x) V_{ta}^\top - \tanh(\mu_t^* A_{ta}^\top x) A_{ta}^\top)(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \\ & = 2 \left(\tanh(\mu_t V_{ta}^\top x) V_{ta}^\top x V_{ta}^\top V_{ta} + \tanh(\mu_t^* A_{ta}^\top x) A_{ta}^\top x A_{ta}^\top A_{ta} \right. \\ & \quad \left. - \tanh(\mu_t^* A_{ta}^\top x) V_{ta}^\top x A_{ta}^\top V_{ta} - \tanh(\mu_t V_{ta}^\top x) A_{ta}^\top x A_{ta}^\top V_{ta} \right). \end{aligned}$$

Then, we have the following bound

$$\begin{aligned} & \mathbb{E}[(\tanh(\mu_t V_{ta}^\top x) V_{ta}^\top - \tanh(\mu_t^* A_{ta}^\top x) A_{ta}^\top)(V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x] \\ & \leq \left(\epsilon(1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta} V_{ta}^\top V_{ta} + \mu_t^* (1 + \mu_t^{*2}) A_{ta}^\top A_{ta} A_{ta}^\top A_{ta} A_{ta}^\top A_{ta} \right. \\ & \quad \left. + \epsilon(1 + \mu_t^{*2}) V_{ta}^\top A_{ta} A_{ta}^\top A_{ta} A_{ta}^\top V_{ta} \right) \triangleq C_2. \end{aligned}$$

$$\begin{aligned}
\text{Bound for } & \left(\frac{1}{\sigma_t^2} - 1\right)^2 \text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)^2 A_{ta}A_{ta}^\top). \\
& \left(\frac{1}{\sigma_t^2} - 1\right)^2 \text{tr}((V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)^2 A_{ta}A_{ta}^\top) \\
& = \left(\frac{1}{\sigma_t^2} - 1\right)^2 [V_{ta}^\top V_{ta}(V_{ta}^\top A_{ta})^2 - 2A_{ta}^\top A_{ta}(V_{ta}^\top A_{ta})^2 + (A_{ta}^\top A_{ta})^3] \triangleq C_3.
\end{aligned}$$

Then, we know that

$$\mathbb{E}[y^\top y] \leq C_1 + 2C_2 + C_3 \triangleq C.$$

For $\forall V_{ta}$, $C = C' + O(\epsilon)$, while

$$\begin{aligned}
C' &= (1 + \mu_t^{*2})\mu_t^{*2}(A_{ta}^\top A_{ta})^3 + 2\mu_t^*(1 + \mu_t^{*2})(A_{ta}^\top A_{ta})^3 + C_3 \\
&= 3(1 + \mu_t^{*2})\mu_t^{*2}(A_{ta}^\top A_{ta})^3 + C_3 < +\infty.
\end{aligned} \tag{10}$$

Then, we obtain the following bound for the gradient of fully fine-tuning method:

$$\mathbb{E}_x \left[\frac{\partial \mathcal{L}_{\text{SM}}^{\text{few-all}}}{\partial \mu_t} \right] \leq 4\epsilon A_{ta}^\top V_{ta} \sqrt{(1 + \mu_t^{*2})V_{ta}^\top V_{ta}} \sqrt{C'} + O(\epsilon^{\frac{3}{2}})$$

When $\epsilon \leq \frac{1}{4V_{ta}^\top A_{ta} \sqrt{C} \sqrt{(1 + \mu_t^{*2})V_{ta}^\top V_{ta}}} \times 10^{-5}$, $\mathbb{E}_x \left[\frac{\partial \mathcal{L}}{\partial \mu_t} \right] \leq 1 \times 10^{-5}$, which indicates that large optimization steps are required in the optimization process.

Under the setting of Example 11, if $V_{ta} = [0.1, 0.1]^\top$, $\mathbb{E}_x \left[\frac{\partial \mathcal{L}_{\text{SM},t}^{\text{few-all}}}{\partial \mu_t} \right] \leq 1 \times 10^{-5}$ when $\epsilon < 0.12$. ■

E THE PROOF FOR GOOD PRETRAINING

In this section, we provide detailed bounds for the symmetric 2-mode case to obtain explicit convergence rates. For a more general derivation covering arbitrary K-mode GMMs and dimensions, please refer to Appendix F. By analyzing the Hessian of score matching objective function for the few-shot phase $\frac{\partial^2 \mathcal{L}_{\text{SM},t}^{\text{few}}}{\partial V_{ta}^2}$, we prove that with a large initialization area, the objective function is strongly convex, which leads to a convergence guarantee.

Since we assume the latent parameter μ^* is perfectly learned by the pretraining phase $\hat{\mu} = \mu^*$, we do not especially distinguish $\hat{\mu}_t$, μ_t^* and μ_t in the proof of this section. We also ignore the subscript t of x when there is no ambiguity. Furthermore, since some results rely on the initialization area, we use the following simple example to show how to satisfy the requirement after providing the theoretical guarantee.

Example Setting

$$t = 2, \mu^* = 4, A_s = [[0.1, 0.1]]^\top \text{ and } A_{ta} = [[0.12, 0.12]]^\top. \tag{11}$$

Recall that the Hessian has the following form

$$2 \underbrace{\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)}_{\text{Squared Term } N} + 2 \underbrace{\left(\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}})}_{\text{Cross Term } M}.$$

First we analyze term $MM^\top$, where M has the form as $aI + bxy^\top$, which will be used in the following lemma.

Lemma E.1. *Let $M = aI + bxy^\top$, $MM^\top$ is semi-positive definite. And it's positive definite if and only if $a = 0$ or $a + bx^\top y = 0$.*

$$\begin{aligned}
& \lambda_{\min}(MM^\top) \\
& = \min \left(a^2, a^2 + abx^\top y + \frac{b^2 \|x\|^2 \|y\|^2}{2} - \frac{b}{2} \sqrt{\|x\|^2 \|y\|^2 (4a^2 + 4abx^\top y + b^2 \|x\|^2 \|y\|^2)} \right)
\end{aligned}$$

Proof. First, $\forall x \in \mathbb{R}^d$, we have

$$\begin{aligned} x^\top MM^\top x &= (M^\top x)^\top (M^\top x) \\ &= \|M^\top x\|_2^2 \geq 0 \end{aligned}$$

Thus, $MM^\top$ is semi-positive definite.

We can also obtain that

$$\begin{aligned} |aI + bxy^\top| &= |aI|(1 + \frac{b}{a}x^\top y) \\ &= a^{n-1}(a + bx^\top y) \end{aligned}$$

Therefore,

$$|MM^\top| = a^{2n-2}(a + bx^\top y)^2 \geq 0,$$

the equality holds if and only if $a = 0$ or $a + bx^\top y = 0$.

We further derive the eigenvalues of $MM^\top$.

Let λ be an eigenvalue of $M^\top M$ with corresponding eigenvector v .

$$(M^\top M)v = \lambda v$$

We can analyze the action of this matrix on two orthogonal subspaces.

Let $S = \text{span}\{x, y\}$. Consider a vector v in the orthogonal complement of S , denoted $S^\perp$. For any such vector $v \neq \mathbf{0}$, we have $x^\top v = 0$ and $y^\top v = 0$.

Let's apply $M^\top M$ to v :

$$\begin{aligned} (M^\top M)v &= (a^2I + ab(xy^\top + yx^\top) + b^2\|x\|^2yy^\top)v \\ &= a^2Iv + ab(x(y^\top v) + y(x^\top v)) + b^2\|x\|^2y(y^\top v) \\ &= a^2v + ab(x(0) + y(0)) + b^2\|x\|^2y(0) \\ &= a^2v \end{aligned}$$

This shows that any vector v orthogonal to both x and y is an eigenvector of $M^\top M$ with the eigenvalue $\lambda = a^2$. The dimension of this subspace, $\dim(S^\perp)$, is at least $n - 2$. Therefore, a^2 is an eigenvalue of $M^\top M$ with a multiplicity of at least $n - 2$.

For the other 2 eigenvalues, we set them μ_1 and μ_2 . We know the determinant of a matrix is the product of its eigenvalues.

$$\det(M^\top M) = (a^2)^{n-2}\mu_1\mu_2$$

We also know that $\det(M^\top M) = \det(M^\top)\det(M) = (\det(M))^2$. The determinant of the original matrix M is $\det(M) = a^{n-1}(a + by^\top x)$. Therefore:

$$\det(M^\top M) = [a^{n-1}(a + by^\top x)]^2 = a^{2n-2}(a + by^\top x)^2$$

Equating the two expressions for the determinant:

$$a^{2n-4}\mu_1\mu_2 = a^{2n-2}(a + by^\top x)^2$$

Solving for the product $\mu_1\mu_2$ (assuming $a \neq 0$):

$$\mu_1\mu_2 = a^2(a + by^\top x)^2$$

The trace of a matrix is the sum of its eigenvalues.

$$\text{tr}(M^\top M) = (n - 2)a^2 + \mu_1 + \mu_2$$

We can also compute the trace directly from the expression for $M^\top M$:

$$\text{tr}(M^\top M) = \text{tr}(a^2 I + ab(xy^\top + yx^\top) + b^2 \|x\|^2 yy^\top)$$

Using the linearity of the trace and the property $\text{tr}(AB) = \text{tr}(BA)$:

$$\begin{aligned} \text{tr}(M^\top M) &= a^2 \text{tr}(I) + ab(\text{tr}(xy^\top) + \text{tr}(yx^\top)) + b^2 \|x\|^2 \text{tr}(yy^\top) \\ &= na^2 + ab(y^\top x + x^\top y) + b^2 \|x\|^2 (y^\top y) \\ &= na^2 + 2ab(y^\top x) + b^2 \|x\|^2 \|y\|^2 \end{aligned}$$

Equating the two expressions for the trace:

$$(n-2)a^2 + \mu_1 + \mu_2 = na^2 + 2ab(y^\top x) + b^2 \|x\|^2 \|y\|^2$$

Solving for the sum $\mu_1 + \mu_2$:

$$\mu_1 + \mu_2 = 2a^2 + 2ab(y^\top x) + b^2 \|x\|^2 \|y\|^2$$

Thus, μ_1 and μ_2 are two solutions of

$$\mu^2 - (2a^2 + 2ab(y^\top x) + b^2 \|x\|^2 \|y\|^2) \mu + a^2(a + by^\top x)^2 = 0$$

We finally obtain that

$$\begin{aligned} \lambda_{\min}(MM^\top) \\ = \min \left(a^2, a^2 + abx^\top y + \frac{b^2 \|x\|^2 \|y\|^2}{2} - \frac{b}{2} \sqrt{\|x\|^2 \|y\|^2 (4a^2 + 4abx^\top y + b^2 \|x\|^2 \|y\|^2)} \right) \end{aligned}$$

■

In the following two lemmas, we provide the bound for the squared term and cross term, respectively.

Lemma E.2. *[Squared Term] Assume Theorem 3.1 and 5.1 holds and the latent parameter $\hat{\mu} = \mu^*$ is learning perfectly. Then, we prove that $N \succeq \alpha I_2$ with $\alpha > 0$ for $\forall t \in [\delta, T]$ (see α in Eq.13).*

Proof. Recall that

$$\begin{aligned} \frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} &= \tanh(\hat{\mu}_t^\top V_{ta}^\top x) \hat{\mu}_t I_2 + \frac{\partial \tanh(\hat{\mu}_t^\top V_{ta}^\top x) \hat{\mu}_t}{\partial V_{ta}} V_{ta}^\top + \left(\frac{1}{\sigma_t^2} - 1 \right) \frac{\partial V_{ta} V_{ta}^\top x}{\partial V_{ta}} \\ &= \tanh(\hat{\mu}_t^\top V_{ta}^\top x) \hat{\mu}_t I_2 + (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x)) \hat{\mu}_t^\top \hat{\mu}_t x V_{ta}^\top + \left(\frac{1}{\sigma_t^2} - 1 \right) (x V_{ta}^\top + V_{ta}^\top x I_2) \\ &= \left(\tanh(\hat{\mu}_t^\top V_{ta}^\top x) \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1 \right) V_{ta}^\top x \right) I_2 \\ &\quad + \left((1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x)) \hat{\mu}_t^\top \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1 \right) \right) x V_{ta}^\top. \end{aligned}$$

Let $p = \tanh(\hat{\mu}_t^\top V_{ta}^\top x) \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1 \right) V_{ta}^\top x$ and $q = (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x)) \hat{\mu}_t^\top \hat{\mu}_t + \frac{1}{\sigma_t^2} - 1$, the squared term can be simplified as:

$$\mathbb{E}_x \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right] = \mathbb{E}_x \left[(pI_2 + qxV_{ta}^\top)(pI_2 + qxV_{ta}^\top)^\top \right].$$

Using lemma E.1, we can obtain that

$$\begin{aligned} & \lambda_{\min}((pI_2 + qxV_{ta}^\top)(pI_2 + qxV_{ta}^\top)^\top) \\ &= \min \left(p^2, p^2 + pqx^\top V_{ta} + \frac{q^2\|x\|^2\|V_{ta}\|^2}{2} - \frac{q}{2}\sqrt{\|x\|^2\|V_{ta}\|^2(4p^2 + 4pqx^\top V_{ta} + q^2\|x\|^2\|V_{ta}\|^2)} \right), \end{aligned}$$

where $p = \tanh(\hat{\mu}_t^\top V_{ta}^\top x)\hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1\right)V_{ta}^\top x$ and $q = (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x))\hat{\mu}_t^\top \hat{\mu}_t + \frac{1}{\sigma_t^2} - 1 > 0$. Moreover, since $q > 0$, we have

$$2pV_{ta}^\top x + q\|x\|^2\|V_{ta}\|^2 \leq \|x\|\|V_{ta}\|\sqrt{4p^2 + 4pqV_{ta}^\top x + q^2\|x\|^2\|V_{ta}\|^2}, \quad (12)$$

the equality holds if and only if $x = kV_{ta}$.

The inequality 12 holds because of the Cauchy-Schwarz Inequality, which can be used through squaring both sides and rearranging the terms.

Thus,

$$\begin{aligned} & pqx^\top V_{ta} + \frac{q^2\|x\|^2\|V_{ta}\|^2}{2} - \frac{q}{2}\sqrt{\|x\|^2\|V_{ta}\|^2(4p^2 + 4pqx^\top V_{ta} + q^2\|x\|^2\|V_{ta}\|^2)} \\ &= \frac{q}{2}(2pV_{ta}^\top x + q\|x\|^2\|V_{ta}\|^2 - \|x\|\|V_{ta}\|\sqrt{4p^2 + 4pqV_{ta}^\top x + q^2\|x\|^2\|V_{ta}\|^2}) \leq 0, \end{aligned}$$

and

$$\begin{aligned} & \lambda_{\min}((pI_2 + qxV_{ta}^\top)(pI_2 + qxV_{ta}^\top)^\top) \\ &= p^2 + pqx^\top V_{ta} + \frac{q^2\|x\|^2\|V_{ta}\|^2}{2} - \frac{q}{2}\sqrt{\|x\|^2\|V_{ta}\|^2(4p^2 + 4pqx^\top V_{ta} + q^2\|x\|^2\|V_{ta}\|^2)}. \end{aligned}$$

After analyzing each term, we can choose $N_1 = \alpha I_2$ with

$$\alpha \triangleq \mathbb{E}_{x \sim q_{ta}} \left[p^2 + pqx^\top V_{ta} + \frac{q^2\|x\|^2\|V_{ta}\|^2}{2} - \frac{q}{2}\sqrt{\|x\|^2\|V_{ta}\|^2(4p^2 + 4pqx^\top V_{ta} + q^2\|x\|^2\|V_{ta}\|^2)} \right], \quad (13)$$

where $p = \tanh(\hat{\mu}_t^\top V_{ta}^\top x)\hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1\right)V_{ta}^\top x$ and $q = (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x))\hat{\mu}_t^\top \hat{\mu}_t + \frac{1}{\sigma_t^2} - 1 > 0$.

Then, we complete our proof. ■

For the cross term, we analyze two situations: the initialization area is around the ground-truth a_{ta} : $|a_{ta} - v_{ta}| \leq \delta_{1,t}$ and initialization area is on the right hand of a_{ta} : $v_{ta} \geq a_{ta} + \delta_{1,t}$. When $v_{ta} \leq a_{ta}$, it is possible for the cross term M to be the negative definite matrix. Hence, we control each element to guarantee the negative influence of the negative definite matrix is small. When $v_{ta} \geq a_{ta}$, the cross term M is semi-positive definite in a large region.

Lemma E.3. [Cross Term] Following setting of Lem. 6.1.

(a) **The** $|a_{ta} - v_{ta}| \leq \delta_{1,t}$ **situation.** For $\forall M(i, j)$, $|M(i, j)| \leq \gamma(\delta_{1,t})$, where $\gamma(\delta_{1,t}) \rightarrow 0$ as $\delta_{1,t} \rightarrow 0$ (see $\gamma(\delta_{1,t})$ in Eq.15).

(b) **The** $v_{ta} \geq a_{ta} + \delta_{1,t}$ **situation.** Let $\delta_{2,t} \triangleq v_{ta} - a_{ta} \geq \delta_{1,t}$ and $M_1 = M - M'$, where M' is SPD. Then, there exists an interval $v_{ta} \in [a_{ta} + \delta_{1,t}, a_{ta} + \delta_{2,t}]$ satisfies:

$$\begin{aligned} & \mathbb{E}[M_1(1, 2)] = \mathbb{E}[M_1(2, 1)] < 0, \\ & \mathbb{E}[M_1(1, 1)] = \mathbb{E}[M_1(2, 2)] > 0, \\ & \mathbb{E}[M_1(1, 1) + M_1(1, 2)] \geq u_1(v_{ta}, t) + u_2(v_{ta}, t), \end{aligned}$$

where $(u_1(v_{ta}, t) + u_2(v_{ta}, t))|_{v_{ta}=a_{ta}+\delta_{1,t}} > 0$, $u_1(\cdot, t)$ increasing and $u_2(\cdot, t)$ decreasing for $v_{ta} \in [a_{ta} + \delta_{1,t}, a_{ta} + \delta_{2,t}]$ (see M' , $u_1(\cdot, t)$ and $u_2(\cdot, t)$ in Eq. 16, 17 and 18).

Proof. We know that the cross term has the following form (in this lemma, we ignore the subscript t of x .)

$$\begin{aligned} & \mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] \\ &= \mathbb{E} [(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t x^\top y I_2 + (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t x y^\top \\ & \quad - 2 \tanh(\mu_t^\top V_{ta}^\top x) (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t \mu_t x^\top y x V_{ta}^\top + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) x^\top y I_2]. \end{aligned} \quad (14)$$

We want to make

$$\begin{aligned} & \mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] + 2 \left(\frac{1}{\sigma_t^2} - 1 \right)^2 (1 + \mu_t^2) A_{ta}^\top A_{ta} V_{ta}^\top V_{ta} I_2 \\ & \quad + \mu_t^2 \left(\frac{1}{\sigma_t^2} - 1 \right) \tanh(\mu_t^2 V_{ta}^\top A_{ta}) V_{ta}^\top A_{ta} I_2 + \mathbb{E}_X [\tanh^2(\mu_t^\top V_{ta}^\top x) \mu_t^2] I_2 \end{aligned}$$

positive definite, where the last three terms come from the above squared term.

In the proof of this lemma, we redefine x :

$$x = [x(1), x(2)]^\top \sim \mathcal{N}(\mu_t A_{ta}, A_{ta} A_{ta}^\top),$$

which indicates that $x(1), x(2) \sim N(\mu_t a_{ta}, a_{ta}^2)$. We also denote by

$$x' \triangleq x(1) + x(2) = [1, 1] \cdot x \sim N(2\mu_t a_{ta}, 4a_{ta}^2).$$

Then, we provide bound for the two situation.

(a) The $|a_{ta} - v_{ta}| \leq \delta_{1,t}$ situation. For any element in the cross term

$$e \in \mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right],$$

we know that

$$\begin{aligned} |\mathbb{E}[e]| &\leq 2|\mathbb{E}[\mu_t^2 (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) x(1) y(1)]| \\ &\quad + 2|\mathbb{E}[\tanh''(\mu_t^\top V_{ta}^\top x) \mu_t^3 (x(1)^2 + x(1)x(2)) y(1)]| + 2\left|\left(\frac{1}{\sigma_t^2} - 1\right) \mathbb{E}[x(1)y(1)]\right| \\ &\leq 2\mu_t^2 \sqrt{\mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) x(1)]^2} \sqrt{\mathbb{E}[y(1)^2]} \\ &\quad + 2\mu_t^3 \sqrt{\mathbb{E}[(\tanh''(\mu_t^\top V_{ta}^\top x) (x(1)^2 + x(1)x(2)))^2]} \sqrt{\mathbb{E}[y(1)^2]} \\ &\quad + 2\left(\frac{1}{\sigma_t^2} - 1\right) \sqrt{\mathbb{E}[x(1)^2]} \sqrt{\mathbb{E}[y(1)^2]}, \end{aligned}$$

where the first inequality follows by the triangle inequality, and the second inequality follows by the Cauchy-Schwarz inequality. Then we give upper bounds on $\mathbb{E}[(\tanh'(\mu_t^\top V_{ta}^\top x))^2 x(1)^2]$, $\mathbb{E}[(\tanh''(\mu_t^\top V_{ta}^\top x) (x(1)^2 + x(1)x(2)))^2]$ and $\mathbb{E}[y^2]$ to obtain a total bound.

(i) Term $\mathbb{E}[(\tanh'(\mu_t^\top V_{ta}^\top x))^2 x(1)^2]$.

With the Cauchy-Schwarz inequality, we know that

$$\begin{aligned} \mathbb{E}[(\tanh'(\mu_t^\top V_{ta}^\top x))^2 x(1)^2] &= \mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 x(1)^2] \\ &\leq \sqrt{\mathbb{E}[\tanh'^4(\mu_t v_{ta} (x(1) + x(2)))]} \sqrt{\mathbb{E}[x(1)^4]}. \end{aligned}$$

For the first component, we know that

$$\begin{aligned}
& \mathbb{E}_{x' \sim N(2\mu_t a_{ta}, 4a_{ta}^2)} [\tanh'^4(\mu_t v_{ta} x')] \\
& \leq \int_0^\infty \tanh'^4(\mu_t v_{ta} x') \exp\left(-\frac{(x' - 2\mu_t a_{ta})^2}{8a_{ta}^2}\right) dx \\
& \stackrel{x' = a_{ta} t}{=} a_{ta} \int_0^\infty \tanh'^4(\mu_t v_{ta} a_{ta} t) \exp\left(-\frac{(t - 2\mu_t)^2}{8}\right) dt \\
& \leq a_{ta} \int_0^\infty \exp(-4\mu_t v_{ta} a_{ta} t) \exp\left(-\frac{(t - 2\mu_t)^2}{8}\right) dt \\
& = a_{ta} \exp(4\mu_t v_{ta} a_{ta} (8\mu_t^2 v_{ta}^2 - 2\mu_t)) \int_0^\infty \exp\left(-\frac{(t + 16\mu_t v_{ta}^2 - 2\mu_t)^2}{8}\right) dt \\
& \leq a_{ta} \exp(4\mu_t^2 v_{ta} a_{ta} (8v_{ta}^2 - 2)).
\end{aligned}$$

Thus,

$$\begin{aligned}
\mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 x(1)^2] & \leq \sqrt{\mathbb{E}[\tanh'^4(\mu_t v_{ta} (x(1) + x(2)))]} \sqrt{\mathbb{E}[x(1)^4]} \\
& \leq \sqrt{a_{ta}} \exp(4\mu_t^2 v_{ta} a_{ta} (4v_{ta}^2 - 1)) \sqrt{3 + 6\mu_t^2 + \mu_t^4 a_{ta}^2} \\
& = \sqrt{3 + 6\mu_t^2 + \mu_t^4 a_{ta}^2} \sqrt{a_{ta}} \exp(4\mu_t^2 v_{ta} a_{ta} (4v_{ta}^2 - 1)),
\end{aligned}$$

where the second inequality follows the fact that $\mathbb{E}[x(1)^4] = (3 + 6\mu_t^2 + \mu_t^4) a_{ta}^4$.

We also know that $0 \leq (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 \leq 1$. As a result, we also can give another bound:

$$\mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 x(1)^2] \leq \mathbb{E}[x(1)^2] = (1 + \mu_t^2) a_{ta}^2$$

Hence, we can obtain that

$$\begin{aligned}
& \mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 x(1)^2] \\
& \leq \min\{\sqrt{3 + 6\mu_t^2 + \mu_t^4 a_{ta}^2} \sqrt{a_{ta}} \exp(4\mu_t^2 v_{ta} a_{ta} (4v_{ta}^2 - 1)), (1 + \mu_t^2) a_{ta}^2\}.
\end{aligned}$$

(ii) Term $\mathbb{E}[(\tanh''(\mu_t^\top V_{ta}^\top x)(x(1)^2 + x(1)x(2)))^2]$.

For this term, we have that

$$\begin{aligned}
\mathbb{E}[(\tanh''(\mu v_{ta} (x(1) + x(2)))(x(1)^2 + x(1)x(2)))^2] & \leq \mathbb{E}[x(1)^2 (x(1) + x(2))^2] \\
& = \mathbb{E}[4x(1)^4] = 4(\mu_t^4 + 6\mu_t^2 + 3) a_{ta}^4,
\end{aligned}$$

where the first inequality holds because $0 \leq \tanh''^2(\mu v_{ta} (x(1) + x(2))) \leq 1$, the second and third equalities hold due to $x(1) = x(2)$ and $\mathbb{E}[x(1)^4] = (\mu_t^4 + 6\mu_t^2 + 3) a_{ta}^4$ respectively.

(iii) Term $\mathbb{E}[y^2]$. Recall that since we assume the latent parameter μ^* is perfectly learned by the pretraining phase $\hat{\mu} = \mu^*$, we do not especially distinguish $\hat{\mu}_t$, μ_t^* and μ_t in the proof of this section.

For y , we have that

$$\begin{aligned}
y & = s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}} \\
& = \mu(V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) - A_{ta} \tanh(\mu_t^\top A_{ta}^\top x)) - (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top)x - \frac{1}{\sigma_t^2} (A_{ta} A_{ta}^\top - V_{ta} V_{ta}^\top)x
\end{aligned}$$

Let $A_{ta} = V_{ta} + \Delta$,

$$\begin{aligned}
& \mathbb{E}[y] \\
& = \mathbb{E}[\mu_t(V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) - A_{ta} \tanh(\mu_t^\top A_{ta}^\top x))] + \mathbb{E}[(1 - \frac{1}{\sigma_t^2})(V_{ta} \Delta^\top + \Delta V_{ta}^\top + \Delta \Delta^\top)x] \\
& \leq \mathbb{E}[\mu_t(V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) - A_{ta} \tanh(\mu_t^\top A_{ta}^\top x))] + (1 - \frac{1}{\sigma_t^2})(V_{ta} \Delta^\top + \Delta V_{ta}^\top + \Delta \Delta^\top) \mu_t A_{ta}.
\end{aligned}$$

We need to give the bound of $\mu_t(V_{ta} \tanh(\mu_t^\top V_{ta}^\top x) - A_{ta} \tanh(\mu_t^\top A_{ta}^\top x))$. Inspired by the Taylor's Theorem, we show $(x + \Delta x) \tanh(x + \Delta x) - x \tanh x$ can be bound by $K \Delta x$, where K will be defined later.

$$\begin{aligned} f(x) &= x \tanh(x) \\ f'(x) &= \tanh(x) + x \cdot \text{sech}^2(x) = \tanh(x) + \frac{4x}{(\exp(x) + \exp(-x))^2}. \end{aligned}$$

For the bound of $f'(x)$, we know that

$$\begin{aligned} |f'(x)| &\leq |\tanh(x)| + \left| \frac{4x}{(\exp(x) + \exp(-x))^2} \right| \\ &\leq \min\{1, |x|\} + \left| \frac{4x}{\exp(2x) + \exp(-2x) + 2} \right| \\ &\leq \min\{1, |x|\} + \min\{|x|, \frac{2}{e}\}, \end{aligned}$$

where the first inequality holds because of the triangle inequality, the second inequality holds because $|\tanh(x)| \leq 1$ and $-x \leq \tanh(x) \leq x$. The third equality holds because

$$\left| \frac{4x}{\exp(2x) + \exp(-2x) + 2} \right| \leq x, \left| \frac{4x}{\exp(2x) + \exp(-2x) + 2} \right| \leq \frac{2}{e}.$$

For $y(1)$ in y , we have that

$$\begin{aligned} |y(1)| &= \left| v_{ta} \tanh(\mu_t(x(1) + x(2))v_{ta}) - (v_{ta} + \delta) \tanh(\mu_t(x(1) + x(2))(v_{ta} + \delta)) + \left(1 - \frac{1}{\sigma_t^2}\right)(2v_{ta}\delta + \delta^2)x(1) \right| \\ &\leq \left| \frac{1}{\mu_t(x(1) + x(2))} \right| \left| (\min\{1, \mu_t(x(1) + x(2))v_{ta}\} + \min\{\mu_t|x(1) + x(2)|v_{ta}, \frac{2}{e}\})\delta \mu_t(x(1) + x(2)) \right| \\ &\quad + \left| \left(1 - \frac{1}{\sigma_t^2}\right)(2v_{ta}\delta + \delta^2)x(1) \right| \\ &\stackrel{\delta \leq 1}{\leq} ((\mu_t(x(1) + x(2))v_{ta} + \min\{\mu_t|x(1) + x(2)|v_{ta}, \frac{2}{e}\})\delta + \left(\frac{1}{\sigma_t^2} - 1\right)(2v_{ta} + 1)\delta x(1)) \\ &= \delta((\min\{1, \mu_t(x(1) + x(2))v_{ta}\} + \min\{\mu_t|x(1) + x(2)|v_{ta}, \frac{2}{e}\} + \left(\frac{1}{\sigma_t^2} - 1\right)(2v_{ta} + 1)x(1)). \end{aligned}$$

Recall that $x' = x(1) + x(2) \sim \mathcal{N}(2\mu_t a_{ta}, 4a_{ta}^2)$. For $\mathbb{E}[y(1)^2]$, we have that

$$\begin{aligned}
\mathbb{E}[y(1)^2] &\leq \mathbb{E}[(\delta((\min\{1, |\mu_t(x(1) + x(2))v_{ta}|\} + \min\{|\mu_t(x(1) + x(2))v_{ta}|, \frac{2}{e}\})) \\
&\quad + \mathbb{E}[(1 - \frac{1}{\sigma_t^2})(2v_{ta} + 1)x(1))^2] \\
&= (\mathbb{E}[(\min\{1, |\mu_t(x(1) + x(2))v_{ta}|\} + \min\{|\mu_t(x(1) + x(2))v_{ta}|, \frac{2}{e}\})^2] \\
&\quad + (1 - \frac{1}{\sigma_t^2})^2(2v_{ta} + 1)^2(1 + \mu_t^2)v_{ta}^2 \mathbb{E}[\min\{1, |\mu_t(x(1) + x(2))v_{ta}|\} \\
&\quad + \min\{|\mu_t(x(1) + x(2))v_{ta}|, \frac{2}{e}\}](1 - \frac{1}{\sigma_t^2})(2v_{ta} + 1)4v_{ta})\delta^2 \\
&\leq (4\mathbb{P}(|\mu_t v_{ta} x'| \geq \frac{2}{e}) + \mathbb{E}_{\mu_t v_{ta} x' < \frac{2}{e}}[4\mu_t^2 v_{ta}^2 x'^2] + (1 - \frac{1}{\sigma_t^2})^2(2v_{ta} + 1)^2(1 + \mu_t^2)v_{ta}^2 \\
&\quad + (\mathbb{E}[2\mu_t v_{ta} x'] + 2\mathbb{P}(|\mu_t v_{ta} x'| \geq \frac{2}{e}))(1 - \frac{1}{\sigma_t^2})(2v_{ta} + 1)4v_{ta})\delta^2 \\
&\leq (4\mathbb{P}(|\mu_t v_{ta} x'| \geq \frac{2}{e}) + 16v_{ta}^2 a_{ta}^2 \mu_t^2 (\mu_t^2 + 1) + (1 - \frac{1}{\sigma_t^2})^2(2v_{ta} + 1)^2(1 + \mu_t^2)v_{ta}^2 \\
&\quad + (4\mu_t^2 v_{ta} a_{ta} + 2\mathbb{P}(|\mu_t v_{ta} x'| \geq \frac{2}{e}))(1 - \frac{1}{\sigma_t^2})(2v_{ta} + 1)4v_{ta})\delta^2 \\
&\triangleq K^2 \delta^2,
\end{aligned}$$

where the first inequality follows by (i) dividing $\mu_t v_{ta} x'$ into two parts $\mu_t v_{ta} x' < 2/e$ and $\mu_t v_{ta} x' \geq 2/e$ (ii) $\min\{1, |\mu_t v_{ta} x'|\} = \min\{2/e, |\mu_t v_{ta} x'|\} = |\mu_t v_{ta} x'|$ when $\mu_t v_{ta} x' < \frac{2}{e}$ and the second inequality follows by $\mathbb{E}_{|\mu_t v_{ta} x'| < \frac{2}{e}}[\mu_t^2 v_{ta}^2 x'^2] \leq \mathbb{E}_x[\mu_t^2 v_{ta}^2 x'^2] = \mu_t^2 v_{ta}^2 a_{ta}^2 (1 + \mu_t^2)$.

For each element in the cross term $e \in \mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right]$, it can be decompose into three term:

$$\begin{aligned}
|\mathbb{E}[e]| &\leq 2 \left| \mathbb{E}[\mu_t^2 (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))x(1)y(1)] \right| + 2 \left| \left(\frac{1}{\sigma_t^2} - 1 \right) \mathbb{E}[x(1)y(1)] \right| \\
&\quad + 2 \left| \mathbb{E}[\tanh''(\mu_t^\top V_{ta}^\top x) \mu_t^3 (x(1)^2 + x(1)x(2))y(1)] \right|.
\end{aligned}$$

For the first term, we have that

$$\begin{aligned}
&2 \left| \mathbb{E}[\mu_t^2 (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))x(1)y(1)] \right| \\
&\leq K\delta \left(2\mu_t^2 \sqrt{\mu_t^4 + 6\mu_t^2 + 3a_{ta}^2 \sqrt{a_{ta}} \exp(4\mu_t^2 v_{ta} a_{ta} (4v_{ta}^2 - 1))} \right).
\end{aligned}$$

For the second term, we have that

$$2 \left| \mathbb{E}[\tanh''(\mu_t^\top V_{ta}^\top x) \mu_t^3 (x(1)^2 + x(1)x(2))y(1)] \right| \leq K\delta \left(2\mu_t^3 \sqrt{4(\mu_t^4 + 6\mu_t^2 + 3)a_{ta}^4} \right).$$

For the third term, we have that

$$2 \left| \left(\frac{1}{\sigma_t^2} - 1 \right) \mathbb{E}[x(1)y(1)] \right| \leq K\delta \left(2 \left(\frac{1}{\sigma_t^2} - 1 \right) \sqrt{1 + \mu_t^2 a_{ta}} \right).$$

Combined with the bound for these three term, we have that

$$\begin{aligned}
&|\mathbb{E}[e]| \\
&\leq 2a_{ta} \mu_t^2 K\delta \times \\
&\quad \left(a_{ta}^{\frac{1}{4}} \sqrt{\mu_t^4 + 6\mu_t^2 + 3} \exp(2\mu_t^2 v_{ta} a_{ta} (4v_{ta}^2 - 1)) + 2\sqrt{\mu_t^4 + 6\mu_t^2 + 3} \mu_t a_{ta} + \left(\frac{1}{\sigma_t^2} - 1 \right) \sqrt{1 + \mu_t^2} \right) \\
&= KC_4 \delta,
\end{aligned} \tag{15}$$

where $\delta \in |\Delta| = |V_{ta} - A_{ta}| \geq 0$.

Now we focus on the Hessian matrix. Let H be the 2×2 Hessian matrix, $\gamma \triangleq KC_4\delta$,

$$\alpha \triangleq \mathbb{E}_{x \sim q_{ta}} \left[p^2 + pqx^\top V_{ta} + \frac{q^2 \|x\|^2 \|V_{ta}\|^2}{2} - \frac{q}{2} \sqrt{\|x\|^2 \|V_{ta}\|^2 (4p^2 + 4pqx^\top V_{ta} + q^2 \|x\|^2 \|V_{ta}\|^2)} \right],$$

where $p = \tanh(\hat{\mu}_t^\top V_{ta}^\top x_t) \hat{\mu}_t + \left(\frac{1}{\sigma_t^2} - 1\right) V_{ta}^\top x_t$ and $q = (1 - \tanh^2(\hat{\mu}_t^\top V_{ta}^\top x_t)) \hat{\mu}_t^\top \hat{\mu}_t + \frac{1}{\sigma_t^2} - 1 > 0$.

As we defined before, we can divide H into two parts H_1 and H_2 :

$$H_1 = \begin{bmatrix} h_1 & 0 \\ 0 & h_1 \end{bmatrix}, h_1 \geq \alpha - \gamma$$

$$H_2 = \begin{bmatrix} h_2 & h_2 \\ h_2 & h_2 \end{bmatrix}, h_2 \geq \alpha' - \gamma,$$

where α and α' is determined in Theorem 6.1. Thus, if $h_1 > 0$ and $h_2 \geq 0$, the Hessian matrix $\frac{\partial^2 \mathcal{L}_{SM,t}^{\text{few}}}{\partial V_{ta}^2}$ is $2(\alpha - \gamma)$ -positive definite.

In our example (Example 11), $\mu_t = 4 \exp(-2)$, $a_{ta} = 0.12$, $\sigma_t = \sqrt{1 - \exp(-4)}$. $\mathbb{P}(|\mu_t v_{ta} x'| \geq \frac{2}{e}) \leq 1 \times 10^{-20} \approx 0$.

Then, we know that when $\delta \leq 0.02$ ($v_{ta} \in [0.1, 0.14]$) $\alpha - \gamma \geq 0$, and

$$h_2 \geq \mathbb{E}_{x(1) \sim \mathcal{N}(\mu_t a_{ta}, a_{ta}^2)} [2(1 - \tanh^2(0.28\mu_t x(1)))^2 \mu_t^2 v_{ta}^2 x(1)^2] - \gamma \geq 0.$$

The $v_{ta} \geq a_{ta} + \delta_{1,t}$ situation. When $v_{ta} \geq a_{ta}$, we will prove that the cross term is semi-positive definite in a large region. If $v_{ta} > a_{ta}$ and $\mu_t = \mu^*$, we can get $x^\top y \geq 0$ and $(1 - \tanh^2(\mu^\top V_{ta}^\top x)) x^\top y \geq 0$:

$$\begin{aligned} x^\top y &= x^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \\ &= x^\top V_{ta} \tanh(\mu_t V_{ta}^\top x) \mu_t - x^\top A_{ta} \tanh(\mu_t^* A_{ta}^\top x) \mu_t^* + \left(\frac{1}{\sigma_t^2} - 1\right) x^\top (V_{ta} V_{ta}^\top - A_{ta} A_{ta}^\top) x \\ &\geq x^\top A_{ta} \tanh(\mu_t A_{ta}^\top x) \mu_t - x^\top A_{ta} \tanh(\mu_t^* A_{ta}^\top x) \mu_t^* + \left(\frac{1}{\sigma_t^2} - 1\right) x^\top (A_{ta} A_{ta}^\top - A_{ta} A_{ta}^\top) x \\ &= 0, \end{aligned}$$

where the inequality holds because $x^\top V_{ta} \tanh(\mu_t V_{ta}^\top x) \mu_t$ is even, monotonically increasing if $V_{ta}^\top x \geq 0$ and $V_{ta}^\top x \geq A_{ta} x$.

Then, we have that

$$\text{tr}(xy^\top) = \text{tr}(y^\top x) = \text{tr}(x^\top y) \geq 0$$

and

$$\text{Rank}(xy^\top) \leq \text{Rank}(x) = 1.$$

We also know that $1 - \tanh^2(\mu_t V_{ta}^\top x) \geq 0$, which indicates $(1 - \tanh^2(\mu_t V_{ta}^\top x)) xy^\top$ is semi-positive definite.

Recall that the cross term has the following form

$$\begin{aligned} M &= \mathbb{E} \left[\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] \\ &= \mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t x^\top y I_2 + (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t x y^\top \\ &\quad - 2 \tanh(\mu_t^\top V_{ta}^\top x) (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t \mu_t x^\top y x V_{ta}^\top + 2 \left(\frac{1}{\sigma_t^2} - 1\right) x^\top y I_2]. \end{aligned}$$

Then, we define the following two matrix: $M = M' + M_1$, where

$$M' = (1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \mu_t^\top \mu_t x y^\top + 2 \left(\frac{1}{\sigma_t^2} - 1\right) x^\top y I_2, \quad (16)$$

and

$$M_1 = (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))\mu_t^2 x^\top y I_2 - 2 \tanh(\mu_t^\top V_{ta}^\top x)\mu_t x V_{ta}^\top (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))\mu_t^2 x^\top y.$$

We know that

$$M_1[1, 1] + M_1[1, 2] = (1 - \tanh^2(\mu_t V_{ta}^\top x))\mu_t^2 x^\top y (1 - 4 \tanh(\mu_t V_{ta}^\top x)\mu_t v_{ta} x(1))$$

$$\begin{aligned} & \mathbb{E}[M_1[1, 1] + M_1[1, 2]] \\ & \geq \mathbb{E}[(1 - \tanh^2(\mu_t V_{ta}^\top x))x^\top y (1 - 4 \tanh(\mu_t V_{ta}^\top x)\mu_t a x(1))] \\ & = \mathbb{E}[(1 - \tanh^2(\mu_t V_{ta}^\top x))x^\top y] - \mathbb{E}[4(1 - \tanh^2(\mu_t V_{ta}^\top x))x^\top y \tanh(\mu_t V_{ta}^\top x(1))\mu_t v_{ta} x(1)]. \end{aligned}$$

Then, we discuss each component in the following part. For the first term, we know that

$$\mathbb{E}[(1 - \tanh^2(\mu_t V_{ta}^\top x))x^\top y] \geq \mathbb{E}[(1 - \tanh^2(\mu_t x(1)))x^\top y] \triangleq u_1(v_{ta}, t). \quad (17)$$

For the second term, we know that

$$\begin{aligned} & -\mathbb{E}[4(1 - \tanh^2(\mu_t V_{ta}^\top x))x^\top y \tanh(\mu_t V_{ta}^\top x(1))\mu_t v_{ta} x(1)] \\ & \geq -\mathbb{E}[4x^\top y \tanh(\mu_t V_{ta}^\top x(1))\mu_t v_{ta} x(1)] \triangleq u_2(v_{ta}, t). \end{aligned} \quad (18)$$

We know that $u_1(v_{ta}, t)$ increases with v_{ta} increasing while $u_2(v_{ta}, t)$ decreases with v_{ta} increasing. We also know that when $v_{ta,t} = a_{ta,t} + \delta_{1,t}$, $u_1(v_{ta}, t) + u_2(v_{ta}, t) > 0$, which indicates there exists an area $v_{ta,t} \in [a_{ta,t} + \delta_{1,t}, a_{ta,t} + \delta_{2,t}]$ that $M_1[1, 1] + M_1[1, 2] \geq 0$.

Thus,

$$\mathbb{E}[M_1[1, 1]] = \mathbb{E}[M_1[2, 2]] > 0, \mathbb{E}[M_1[1, 2]] = \mathbb{E}[M_1[2, 1]] < 0,$$

and

$$\begin{aligned} |\mathbb{E}[M_1]| &= (\mathbb{E}[M_1[1, 1]])^2 - (\mathbb{E}[M_1[1, 2]])^2 \\ &= (\mathbb{E}[M_1[1, 1]] + \mathbb{E}[M_1[1, 2]])(\mathbb{E}[M_1[1, 1]] - \mathbb{E}[M_1[1, 2]]) > 0. \end{aligned}$$

Then we know that

$$\mathbb{E}_x[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))\mu_t^2 x^\top y I_2 - 2 \tanh(\mu_t^\top V_{ta}^\top x)\mu_t x V_{ta}^\top (1 - \tanh^2(\mu_t^\top V_{ta}^\top x))\mu_t^2 x^\top y]$$

is semi-positive definite. Then, the proof is finished.

To make a clearer discussion, we use the setting of Example 11 to show the interval of $[a_{ta} + \delta_{1,t}, a_{ta} + \delta_{2,t}]$.

$$v_{ta} \in [0.14, 0.25]$$

$$\begin{aligned} u_1(0.14) &\approx 0.00023 > 0.0002 \\ u_2(0.28) &\leq 4 \times 10^{-5} < f(0.14) \end{aligned}$$

$$v_{ta} \in [0.25, 0.4]$$

$$\begin{aligned} u_1(0.25) &\approx 0.0021 > 0.002 \\ u_2(0.4) &\leq 1.4 \times 10^{-4} < f(0.25) \end{aligned}$$

$$v_{ta} \in [0.4, 0.5]$$

$$\begin{aligned} u_1(0.4) &\approx 0.0064 > 0.006 \\ u_2(0.5) &\leq 0.00034 < f(0.4) \end{aligned}$$

Hence, we can have $\mathbb{E}[M_1(1, 1) + M_1(1, 2)] > 0$ when $v_{ta} \in [0.14, 0.5]$. ■

Before proving our convergence guarantee, we first previous convergence lemma.

Lemma E.4 (Convergence Lemma). *Let ϕ be locally μ -strongly convex and L_m -smooth, if $\eta_t = \eta = \frac{2}{\mu + L_m}$, $\kappa = \frac{L_m}{\mu}$, and $x^* \in \arg\min_{x \in \mathcal{X}} \phi(x)$, then*

$$\|x^t - x^*\|_2 \leq \left(\frac{\kappa - 1}{\kappa + 1} \right)^t \|x^{(0)} - x^*\|_2.$$

After that, we provide our convergence guarantee for few-shot diffusion models with a great pretraining.

Theorem 6.5. *Assume Theorem 3.1, 5.1, $\hat{\mu} = \mu^*$ and $\delta_{1,t}, \delta_{2,t}$ satisfy **Condition 1**. Considering score matching function $\mathcal{L}$. When $v_{ta}^{(0)} \in \{[a_{ta} - \delta_{1,t}, a_{ta} + \delta_{2,t}] \cup [-a_{ta} - \delta_{2,t}, -a_{ta} + \delta_{1,t}]\}$, using gradient descent with learning rate $\eta = 1/(2\alpha + \zeta)$, we have*

$$\begin{aligned} & \left\| V_{ta}^{(k)} V_{ta}^{(k)\top} - A_{ta} A_{ta}^\top \right\|_F \\ & \leq \left(\frac{\kappa - 1}{\kappa + 1} \right)^k (2a_{ta} + \delta_{2,t}) |v_{ta}^{(0)} - a_{ta}|, \end{aligned}$$

where $\kappa = (\alpha + \gamma + \zeta)/(\alpha - \gamma)$.

Proof. First we prove that there exists $L_m > 0$, such that the objective function is L_m -smooth. In this work, we take the maximum eigenvalue of the hessian matrix to be L_m .

$$\mathbb{E} \left[\frac{\partial^2 \mathcal{L}_{SM}^{\text{few}}}{\partial V_{ta}^2} \right] = 2\mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2} \right)^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}}) \right] + 2\mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right]$$

Based on our analysis of the hessian matrix, we can divide the matrix into two parts: $\begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_1 \end{bmatrix}$ and $\begin{bmatrix} \lambda_2 & \lambda_2 \\ \lambda_2 & \lambda_2 \end{bmatrix}$.

We first analyze the property of $\begin{bmatrix} \lambda_1 + \lambda_2 & \lambda_2 \\ \lambda_2 & \lambda_1 + \lambda_2 \end{bmatrix}$, and then give the bound of λ_1 and λ_2 .

$$\left| \lambda I_2 - \begin{bmatrix} \lambda_1 + \lambda_2 & \lambda_2 \\ \lambda_2 & \lambda_1 + \lambda_2 \end{bmatrix} \right| = 0 \Rightarrow (\lambda - \lambda_1)(\lambda - \lambda_1 - \lambda_2) = 0,$$

which indicates $\lambda = \lambda_1$ or $\lambda = \lambda_1 + \lambda_2$. Thus, if $\lambda_1 > 0$, we can choose $L_m = \lambda_1 + |\lambda_2|$

According to our analysis on before, $\forall e \in \mathbb{E}[(\frac{\partial^2 s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}^2})^\top (s_{\hat{\mu}, V_{ta}} - s_{\mu^*, A_{ta}})]$, $|e| \leq v$

Next we analyze $\mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right]$ and have the following form:

$$\begin{aligned} & \mathbb{E} \left[\left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right)^\top \left(\frac{\partial s_{\hat{\mu}, V_{ta}}}{\partial V_{ta}} \right) \right] \\ & = \mathbb{E}[\tanh^2(\mu_t^\top V_{ta}^\top x) \mu_t^\top \mu_t I_2] + \mathbb{E}[(1 - \tanh^2(\mu_t^\top V_{ta}^\top x))^2 \mu_t^\top \mu_t V_{ta} x^\top x V_{ta}^\top] \\ & + 2\left(\frac{1}{\sigma_t^2} - 1\right)^2 ((1 + \mu_t^2) A_{ta} A_{ta}^\top V_{ta} V_{ta}^\top + (1 + \mu_t^2) V_{ta}^\top A_{ta} A_{ta}^\top V_{ta}) \\ & + \mathbb{E}[2(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \tanh(\mu_t^\top V_{ta}^\top x) \mu_t^\top \mu_t V_{ta} x^\top] \\ & + 2\mathbb{E}[(\frac{1}{\sigma_t^2} - 1) \tanh(\mu_t^\top V_{ta}^\top x) \mu_t (x V_{ta}^\top + V_{ta}^\top x I_2)] \\ & + \mathbb{E}(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \left(\frac{1}{\sigma_t^2} - 1 \right) \mu_t^\top \mu_t (V_{ta} x^\top x V_{ta}^\top + V_{ta} x^\top V_{ta}^\top x) \\ & + \mathbb{E}(1 - \tanh^2(\mu_t^\top V_{ta}^\top x)) \left(\frac{1}{\sigma_t^2} - 1 \right) \mu_t^\top \mu_t (x^\top V_{ta} x V_{ta}^\top + x^\top V_{ta} V_{ta}^\top x) \\ & = \begin{bmatrix} \alpha & 0 \\ 0 & \alpha \end{bmatrix} + \begin{bmatrix} \zeta & \zeta \\ \zeta & \zeta \end{bmatrix}, \end{aligned}$$

where

$$\begin{aligned} \zeta = \mathbb{E}_{x(1) \sim \mathcal{N}(\mu_t a_{ta}, a_{ta}^2)} & \left[2(1 - \tanh^2(2\mu_t v_{ta} x(1)))^2 \mu_t^2 v_{ta}^2 x(1)^2 \right. \\ & + 2(1 - \tanh^2(2\mu_t v_{ta} x(1))) \tanh(2\mu_t v_{ta} x(1)) \mu_t^3 v_{ta} x(1) \\ & + 2 \left(\frac{1}{\sigma_t^2} - 1 \right) \tanh(2\mu_t v_{ta} x(1)) \mu_t a x(1) \\ & \left. + \left(\frac{1}{\sigma_t^2} - 1 \right) (1 - \tanh^2(2\mu_t v_{ta} x(1))) 6\mu_t^2 v_{ta}^2 x(1)^2 + \left(\frac{1}{\sigma_t^2} - 1 \right)^2 4v_{ta}^4 x(1)^4 \right]. \end{aligned} \quad (19)$$

For the ζ , we have the following bound:

$$\begin{aligned} \zeta & \leq \mathbb{E} \left[2\mu_t^2 v_{ta}^2 x(1)^2 + \left(\frac{1}{\sigma_t^2} - 1 \right)^2 4v_{ta}^4 x(1)^4 + 4\mu_t^4 v_{ta}^2 x(1)^2 \right] \\ & + 4\mathbb{E} \left[\left(\frac{1}{\sigma_t^2} - 1 \right) \mu_t^4 v_{ta}^2 x(1)^2 + 6 \left(\frac{1}{\sigma_t^2} - 1 \right) \mu_t^2 v_{ta}^2 x(1)^2 \right] \\ & = \mu_t^2 v_{ta}^2 a_{ta}^2 (1 + \mu_t^2) [2 + 4\mu_t^2 + 4 \left(\frac{1}{\sigma_t^2} - 1 \right) \mu_t^2 + 6 \left(\frac{1}{\sigma_t^2} - 1 \right)] \\ & + 4 \left(\frac{1}{\sigma_t^2} - 1 \right)^2 (\mu_t^4 + 6\mu_t^2 + 3) v_{ta}^4 a_{ta}^4. \end{aligned}$$

Thus, we can take $L_m = 2(\alpha + \gamma + \zeta)$. Let $\kappa = \frac{L_m}{\alpha - \gamma}$, $\eta = \frac{2}{2(\alpha - \gamma) + 2(\alpha + \gamma + \zeta)} = \frac{1}{2\alpha + \zeta}$ and $A_{ta} \in \operatorname{argmin}_{V_{ta} \in \mathcal{Q}(\hat{\mu})} L_{SM,t}^{\text{few}}$, then

$$\left\| V_{ta}^{(k)} V_{ta}^{(k)\top} - A_{ta} A_{ta}^\top \right\|_F \leq \left(\frac{\kappa - 1}{\kappa + 1} \right)^k (2a_{ta} + \delta_{2,t}) |v_{ta}^{(0)} - a_{ta}|.$$

■

F DERIVATION AND ANALYSIS OF K-MODE GMM

In this section, we derive the score function and its Jacobian matrix for the K-mode Gaussian Mixture Model (GMM) with latent dimension $d = 1$. We strictly follow the notation established in the previous sections. By analyzing the structure of the Jacobian, we derive the Squared Term (N) of the Hessian to prove strong convexity under good initialization.

F.1 SCORE FUNCTION AND JACOBIAN DERIVATION

For the K-mode GMM, the score function $s_\theta(x)$ is defined as:

$$s_\theta(x) = V_{ta} \bar{\mu} + \left(\frac{1}{\sigma_t^2} - 1 \right) V_{ta} V_{ta}^\top x - \frac{1}{\sigma_t^2} x, \quad (20)$$

where $z = V_{ta}^\top x$ is the latent projection, and $\bar{\mu} = \sum_{i=1}^K \omega_{i,t} \hat{\mu}_{i,t}$ is the weighted mean of the modes. The weights are given by $\omega_{i,t} = \frac{e_{i,t}}{\sum_{k=1}^K e_{k,t}}$, with unnormalized probabilities $e_{i,t} = \exp(-\frac{1}{2}(z - \hat{\mu}_{i,t})^2)$.

We seek the Jacobian matrix $J = \frac{\partial s_\theta}{\partial V_{ta}} \in \mathbb{R}^{D \times D}$.

Lemma F.1. *The Jacobian matrix J of the K-mode GMM score function admits a rank-1 perturbed identity form:*

$$J = \alpha I_D + \beta V_{ta} x^\top,$$

where the scalar coefficients are defined as:

$$\begin{aligned}\alpha &= \bar{\mu} + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta}^\top x), \\ \beta &= \sigma_{var}^2 + \left(\frac{1}{\sigma_t^2} - 1 \right) .\end{aligned}$$

Here, $\sigma_{var}^2 = \sum_{i=1}^K \omega_{i,t} (\hat{\mu}_{i,t} - \bar{\mu})^2$ represents the weighted variance of the modes in the latent space.

Proof. We compute the derivatives of the non-linear term $(V_{ta}\bar{\mu})$ and the linear term $(V_{ta}V_{ta}^\top x)$ separately.

(i) Derivative of the Non-linear Term: Using the product rule, the derivative of $V_{ta}\bar{\mu}$ with respect to V_{ta} is:

$$\frac{\partial(V_{ta}\bar{\mu})}{\partial V_{ta}} = \bar{\mu}I_D + V_{ta} \left(\frac{\partial \bar{\mu}}{\partial V_{ta}} \right)^\top . \quad (21)$$

To find $\frac{\partial \bar{\mu}}{\partial V_{ta}}$, we differentiate the weighted sum. Assuming $\hat{\mu}_{i,t}$ are fixed parameters relative to the projection:

$$\frac{\partial \bar{\mu}}{\partial V_{ta}} = \sum_{i=1}^K \hat{\mu}_{i,t} \frac{\partial \omega_{i,t}}{\partial V_{ta}} .$$

We first compute the derivative of the unnormalized probability $e_{i,t}$ using the chain rule:

$$\begin{aligned}\frac{\partial e_{i,t}}{\partial V_{ta}} &= e_{i,t} \cdot \frac{\partial}{\partial V_{ta}} \left[-\frac{1}{2} (V_{ta}^\top x - \hat{\mu}_{i,t})^2 \right] \\ &= -e_{i,t} (z - \hat{\mu}_{i,t}) x .\end{aligned}$$

Applying the quotient rule to the softmax function $\omega_{i,t}$:

$$\begin{aligned}\frac{\partial \omega_{i,t}}{\partial V_{ta}} &= \frac{1}{\sum e_{k,t}} \frac{\partial e_{i,t}}{\partial V_{ta}} - \frac{e_{i,t}}{(\sum e_{k,t})^2} \sum_{k=1}^K \frac{\partial e_{k,t}}{\partial V_{ta}} \\ &= \omega_{i,t} \left[-(z - \hat{\mu}_{i,t}) - \sum_{k=1}^K \omega_{k,t} (-(z - \hat{\mu}_{k,t})) \right] x \\ &= \omega_{i,t} [(\hat{\mu}_{i,t} - z) - (\bar{\mu} - z)] x \\ &= \omega_{i,t} (\hat{\mu}_{i,t} - \bar{\mu}) x .\end{aligned}$$

Substituting this back into the expression for $\frac{\partial \bar{\mu}}{\partial V_{ta}}$:

$$\frac{\partial \bar{\mu}}{\partial V_{ta}} = \left[\sum_{i=1}^K \omega_{i,t} \hat{\mu}_{i,t} (\hat{\mu}_{i,t} - \bar{\mu}) \right] x = \sigma_{var}^2 x .$$

Thus, Eq. (21) becomes:

$$\frac{\partial(V_{ta}\bar{\mu})}{\partial V_{ta}} = \bar{\mu}I_D + \sigma_{var}^2 V_{ta} x^\top . \quad (22)$$

(ii) Derivative of the Linear Term: We differentiate $(\frac{1}{\sigma_t^2} - 1)V_{ta}(V_{ta}^\top x)$. Noting that $V_{ta}^\top x$ is a scalar:

$$\begin{aligned}\frac{\partial}{\partial V_{ta}} \left[\left(\frac{1}{\sigma_t^2} - 1 \right) V_{ta} (V_{ta}^\top x) \right] &= \left(\frac{1}{\sigma_t^2} - 1 \right) \left[(V_{ta}^\top x) I_D + V_{ta} \frac{\partial (V_{ta}^\top x)}{\partial V_{ta}} \right] \\ &= \left(\frac{1}{\sigma_t^2} - 1 \right) \left[(V_{ta}^\top x) I_D + V_{ta} x^\top \right] .\end{aligned} \quad (23)$$

(iii) **Combined Jacobian:** Summing the derivatives from Eq. (22) and Eq. (23):

$$\begin{aligned}
J &= [\bar{\mu}I_D + \sigma_{var}^2 V_{ta} x^\top] + \left(\frac{1}{\sigma_t^2} - 1 \right) [(V_{ta}^\top x)I_D + V_{ta} x^\top] \\
&= \left[\bar{\mu} + \left(\frac{1}{\sigma_t^2} - 1 \right) (V_{ta}^\top x) \right] I_D + \left[\sigma_{var}^2 + \left(\frac{1}{\sigma_t^2} - 1 \right) \right] V_{ta} x^\top \\
&= \alpha I_D + \beta V_{ta} x^\top,
\end{aligned}$$

matching the form in Lemma E.1. ■

F.2 ANALYSIS OF THE SQUARED TERM

The Squared Term of the Hessian is approximated by $N = JJ^\top$. Since J is in the form $\alpha I + \beta uv^\top$ (where $u = V_{ta}, v = x$), we apply **Lemma F.1** to determine its spectral properties.

The eigenvalues of N determine the local convexity. Using the closed-form solution for the eigenvalues of a rank-1 perturbed identity matrix, the minimum eigenvalue is given by:

$$\lambda_{\min}(N) = \min \left(\alpha^2, \quad \alpha^2 + \alpha\beta(V_{ta}^\top x) + \frac{\beta^2 \|V_{ta}\|^2 \|x\|^2}{2} - \frac{|\beta|}{2} \sqrt{\mathcal{D}} \right), \quad (24)$$

where the discriminant $\mathcal{D}$ is:

$$\mathcal{D} = \|V_{ta}\|^2 \|x\|^2 (4\alpha^2 + 4\alpha\beta(V_{ta}^\top x) + \beta^2 \|V_{ta}\|^2 \|x\|^2).$$

This analytical form allows us to bound the smallest eigenvalue away from zero, provided that the initialization is sufficiently close to the ground truth (ensuring $\alpha \neq 0$).

F.3 HESSIAN ANALYSIS

Recall that the Hessian consists of the Squared Term N and the Cross Term M_{cross} :

$$H = \underbrace{2 \mathbb{E} [J^\top J]}_N + \underbrace{2 \mathbb{E} \left[\frac{\partial^2 s_\theta}{\partial V_{ta}^2} (s_\theta - s_{target}) \right]}_{M_{cross}}.$$

F.3.1 Analysis of the Squared Term

Using Lemma F.1, the matrix product $JJ^\top$ (before expectation) is:

$$\begin{aligned}
JJ^\top &= (\alpha I + \beta V_{ta} x^\top)(\alpha I + \beta x V_{ta}^\top) \\
&= \alpha^2 I + \alpha\beta(V_{ta} x^\top + x V_{ta}^\top) + \beta^2 \|V_{ta}\|^2 x x^\top.
\end{aligned}$$

Applying Lemma E.1, we can lower bound the eigenvalues of this term. We define the positive definite constant for the squared term as λ_{sq} :

$$\lambda_{sq} = \mathbb{E}_{x \sim q_{ta}} \left[\alpha^2 + \alpha\beta(x^\top V_{ta}) + \frac{\beta^2 \|x\|^2 \|V_{ta}\|^2}{2} - \frac{|\beta|}{2} \sqrt{\mathcal{D}} \right], \quad (25)$$

where $\mathcal{D}$ is the discriminant defined in Lemma E.1. Since the data is generated from a GMM, the moments of x are finite, ensuring $\lambda_{sq} > 0$.

F.3.2 Bounding the Cross Term

We now analyze the residual term $y(x) = s_\theta(x) - s_{target}(x)$ induced by the parameter error $\Delta = V_{ta} - A_{ta}$.

Lemma F.2. *The residual norm is strictly bounded by the parameter error:*

$$\|y(x)\|_2 \leq \mathcal{K}(x)\|\Delta\|_2,$$

where $\mathcal{K}(x)$ is a state-dependent coefficient affine in $\|x\|_2$.

Proof. Decomposing the residual into linear and non-linear parts:

$$y(x) = \underbrace{\left(\frac{1}{\sigma_t^2} - 1\right)(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)x}_{\text{Term A}} + \underbrace{(V_{ta}\bar{\mu}(V_{ta}^\top x) - A_{ta}\bar{\mu}(A_{ta}^\top x))}_{\text{Term B}}.$$

For Term A, utilizing $VV^\top - AA^\top \approx A\Delta^\top + \Delta A^\top$:

$$\left\| \left(\frac{1}{\sigma_t^2} - 1\right)(V_{ta}V_{ta}^\top - A_{ta}A_{ta}^\top)x \right\|_2 \leq \left| \frac{1}{\sigma_t^2} - 1 \right| (2\|A_{ta}\| + \|\Delta\|)\|x\|\|\Delta\|.$$

For Term B, we utilize the fact that the derivative of the mean function is the variance σ_{var}^2 , which is bounded by a constant $L_{\bar{\mu}}$. Thus $\bar{\mu}$ is $L_{\bar{\mu}}$ -Lipschitz continuous:

$$\|(V_{ta}\bar{\mu}(V_{ta}^\top x) - A_{ta}\bar{\mu}(A_{ta}^\top x))\|_2 \leq \mu_{\max}\|\Delta\| + \|V_{ta}\|L_{\bar{\mu}}\|\Delta\|\|x\|.$$

Summing the bounds for Term A and Term B, we obtain:

$$\begin{aligned} \|y(x)\|_2 &\leq \underbrace{\mu_{\max}}_{\kappa_0}\|\Delta\| + \underbrace{\left[\left| \frac{1}{\sigma_t^2} - 1 \right| (2\|A_{ta}\| + \|\Delta\|) + L_{\bar{\mu}}\|V_{ta}\| \right]}_{\kappa_1}\|x\|\|\Delta\| \\ &= (\kappa_0 + \kappa_1\|x\|_2)\|\Delta\|_2. \end{aligned}$$

Thus, explicitly, $\mathcal{K}(x) = \kappa_0 + \kappa_1\|x\|_2$, which is affine in $\|x\|_2$. ■

Using Lemma F.2, the residual term $y(x)$ is bounded. To bound the full Cross Term M_{cross} , we must also bound the Hessian of the score function with respect to the parameters, denoted as $\mathcal{H}_{score}(x) = \frac{\partial^2 s_\theta}{\partial V_{ta}^2}$.

Analysis of Eq. (20) allows us to derive the explicit bounds for the Hessian tensor $\mathcal{H}_{score}(x)$. The Hessian consists of a constant component derived from the linear diffusion term and a dynamic component derived from the GMM moments. Specifically, differentiating the linear term yields a bound proportional to $2\left|\frac{1}{\sigma_t^2} - 1\right|\|x\|$, while differentiating the variance term σ_{var}^2 in the Jacobian introduces the third central moment, which scales with $\|x\|^2$.

Thus, we explicitly define the coefficients c_1 and c_2 :

$$\begin{aligned} c_1 &= 2 \left| \frac{1}{\sigma_t^2} - 1 \right| + \mu_{\max}^2, \\ c_2 &= \mu_{\max}\|A_{ta}\|_2, \end{aligned}$$

where $\mu_{\max} = \max_i \|\hat{\mu}_{i,t}\|$ is the maximum norm of the mode centers, representing the bound on the variance and skewness of the latent GMM.

Substituting these into the expectation:

$$\begin{aligned} \|M_{cross}\|_2 &\leq \mathbb{E}[(c_1\|x\| + c_2\|x\|^2)\mathcal{K}(x)\|\Delta\|] \\ &= \|\Delta\| (c_1\kappa_0m_1 + (c_1\kappa_1 + c_2\kappa_0)m_2 + c_2\kappa_1m_3). \end{aligned}$$

Let

$$C_{cross} = c_1\kappa_0m_1 + (c_1\kappa_1 + c_2\kappa_0)m_2 + c_2\kappa_1m_3. \quad (26)$$

This explicit form confirms that data with larger noise variance, larger mode separation, or heavier tails (large moments m_k) results in a larger C_{cross} , thereby requiring a tighter initialization bound.

F.4 STRONG CONVEXITY AND CONVERGENCE

Theorem F.3 (Strong Convexity under K-mode Initialization). *Assume the pretraining phase yields an initialization $V_{ta}^{(0)}$ such that the parameter error satisfies $\|\Delta\|_{init} < \frac{\lambda_{sq}}{C_{cross}}$. Then, the Hessian is strictly positive definite:*

$$\lambda_{\min}(H) \geq 2\lambda_{sq} - 2C_{cross}\|\Delta\|_{init} > 0,$$

recall that $\Delta = V_{ta} - A_{ta}$.

Consequently, the optimization objective is locally strongly convex, guaranteeing linear convergence to the ground truth A_{ta} .

Proof. The Hessian matrix is given by $H = 2N - 2M_{cross}$. To prove strong convexity, we examine the minimum eigenvalue of H . Using the property that $\lambda_{\min}(A - B) \geq \lambda_{\min}(A) - \|B\|_2$ for symmetric matrices, we have:

$$\begin{aligned} \lambda_{\min}(H) &\geq \lambda_{\min}(2N) - \|2M_{cross}\|_2 \\ &= 2\lambda_{\min}(N) - 2\|M_{cross}\|_2. \end{aligned}$$

We now substitute the bounds established in the previous sections:

(i) Squared Term

For the Squared Term, analysis in Section F ensures that near the ground truth, N is positive definite with $\lambda_{\min}(N) \geq \lambda_{sq}$, where

$$\lambda_{sq} \triangleq \mathbb{E}_{x \sim q_{ta}} \left[\alpha^2 + \alpha\beta(x^\top V_{ta}) + \frac{\beta^2 \|x\|^2 \|V_{ta}\|^2}{2} - \frac{|\beta|}{2} \sqrt{\mathcal{D}} \right]. \quad (27)$$

(ii) The Cross Term

For the Cross Term, applying Lemma F.2 and the subsequent moment analysis yields the bound $\|M_{cross}\|_2 \leq C_{cross}\|\Delta\|$. Substituting these terms back into the eigenvalue inequality:

$$\lambda_{\min}(H) \geq 2\lambda_{sq} - 2C_{cross}\|\Delta\|.$$

For the landscape to be locally strongly convex ($\lambda_{\min}(H) > 0$), we require:

$$2(\lambda_{sq} - C_{cross}\|\Delta\|) > 0 \implies \|\Delta\| < \frac{\lambda_{sq}}{C_{cross}}.$$

Under this condition, the objective is μ -strongly convex with $\mu = 2(\lambda_{sq} - C_{cross}\|\Delta\|)$, guaranteeing linear convergence. ■